

ACM/CS 114

Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Winter 2010

Dense matrix problems

- ▶ we'll take a look at
 - ▶ inner and outer products of two vector
 - ▶ matrix-vector and matrix-matrix multiplication
 - ▶ LU factorization and Cholesky decomposition
 - ▶ QR factorization
 - ▶ computing eigenvalues and eigenvectors
 - ▶ fast Fourier transforms
- ▶ when solving a problem of size n on p processors, we will assume
 - ▶ that p , and occasionally \sqrt{p} divides n
 - ▶ that p is a perfect square, when forming two-dimensional process grids
 - ▶ matrices are $n \times n$ – square, not rectangular
 - ▶ we are memory constrained and data replication must be minimized
- ▶ these problems have been studied extensively and form the core of scientific computing on parallel machines
 - ▶ excellent implementations available
 - ▶ interest has been revived due to the expected disruption by multi-core architectures

Vector inner product

- ▶ the inner product of two n -vectors x, y is given by

$$x^T y = \sum_{i=1}^n x_i y_i \quad (1)$$

which requires n multiplications and $n - 1$ additions

- ▶ parallelization strategy:
 - ▶ n fine grain tasks, numbered $i = 1, \dots, n$, that store x_i and y_i , and compute $x_i y_i$
 - ▶ communication is a sum reduction over n fine grain tasks
 - ▶ coarsening is achieved by coalescing n/p tasks together, assuming that each process can accommodate the data storage requirements
 - ▶ and mapping each coarse grain task to a process

Vector outer product

- ▶ the outer product of two n -vectors x and y is the $n \times n$ matrix A given by

$$A_{ij} = x_i y_j \quad (2)$$

which requires n^2 multiplications

- ▶ parallelization strategies are determined by the storage requirements
 - ▶ build a two-dimensional grid of n^2 fine grain tasks numbered (i, j) , with $i, j = 1, \dots, n$; each one computes $x_i y_j$
 - ▶ assuming no data replication is allowed
 - ▶ let task $(i, 1)$ store x_i and task $(1, i)$ store y_i
 - ▶ or, let task (i, i) store both x_i and y_i
 - ▶ either way, the task that owns each element must broadcast it to the other tasks: x_i along the i^{th} task row, y_j along the j^{th} task column
 - ▶ coarsening to p tasks can be accomplished by
 - ▶ combining n/p rows or columns
 - ▶ forming $(n/\sqrt{p}) \times (n/\sqrt{p})$ grid of fine grain tasks
 - ▶ and each coarse grain task can be assigned to a process
- ▶ either way, naïve broadcasting of the components of x and y would require as much total memory as replication
 - ▶ storage can be reduced by circulating portions of x and y through the tasks, with each task using the available portion and passing it on

The product of a matrix with a vector

- ▶ given an $n \times n$ matrix A and an n -vector x , the matrix vector product yields an n -vector y whose components are given by

$$y_i = \sum_{j=1}^n A_{ij}x_j \quad (3)$$

requiring a total of n^2 multiply-add operations

- ▶ once again, the parallelization strategy is determined by how the data is distributed among fine grain tasks
 - ▶ build a two-dimensional grid of n^2 fine grain tasks numbered (i,j) , with $i, j = 1, \dots, n$; each one computes $A_{ij}x_j$
 - ▶ task (i,j) has A_{ij} , but if no data replication is allowed
 - ▶ let task $(i, 1)$ store x_i and task $(1, i)$ store y_i
 - ▶ or, let task (i, i) store both x_i and y_i
 - ▶ the task that owns x_j must broadcast it along the j^{th} task row, and y_i is formed by sum reduction along the i^{th} task column
 - ▶ coarsening into p tasks can be accomplished by combining n/p rows/columns, or by forming $(n/\sqrt{p}) \times (n/\sqrt{p})$ blocks
 - ▶ and each coarse grain task can be assigned to a process

Coarsening along rows or columns

- ▶ for one-dimensional coarsening into n/p task rows
 - ▶ if x is stored in one task, it must be broadcast to all others
 - ▶ if x is distributed among tasks, with n/p components per task, then multiple broadcasts are required
 - ▶ each task computes the inner product of its n/p rows of A with the entire x to produce n/p components of y
- ▶ for one-dimensional coarsening into n/p task columns
 - ▶ n/p components of x are distributed among the tasks
 - ▶ each task computes the linear combination of its n/p columns with coefficients from its copy of x
 - ▶ since the right parts of x are already available, no communication is required
 - ▶ y is generated by a sum reduction across tasks
- ▶ these two are *duals* of each other
 - ▶ row coarsening begins with broadcast, followed by communication-free inner products
 - ▶ column coarsening begins with communication-free linear combinations, follows by a reduction

Two dimensional coarsening

- ▶ for two dimensional coarsening, we form $(n/\sqrt{p}) \times (n/\sqrt{p})$ blocks of fine grain task
 - ▶ each one holding a $(n/\sqrt{p}) \times (n/\sqrt{p})$ block of A
 - ▶ with components of x distributed either across one task row, or along the diagonal, n/p components per task
- ▶ the algorithm combines the features of row/column coarsening
 - ▶ components of x are broadcast along task columns
 - ▶ each task performs n^2/p multiplications locally and sums n/\sqrt{p} sets of products
 - ▶ sum reductions along task rows produce the components of y by combining the component products

Matrix multiplication

- ▶ the product of two $n \times n$ matrices A and B is an $n \times n$ matrix C given by

$$C_{ij} = \sum_{k=1}^n A_{ik} B_{kj} \quad (4)$$

where each one of n^2 entries requires n multiply-adds for a total of n^3 operations

- ▶ matrix multiplication can be viewed as
 - ▶ n^2 inner products
 - ▶ the sum of n outer products
 - ▶ n matrix vector products
- ▶ each one produces a parallel algorithm for matrix multiplication
- ▶ but we'll explore a direct solution instead

Partitioning and communication patterns

- ▶ we build a three dimensional array of n^3 fine grain tasks
 - ▶ with $i, j, k = 1, \dots, n$, let task (i, j, k) be responsible for computing the product $A_{ij}B_{jk}$
 - ▶ assuming no data replication, we have to distribute the data for A and B among $2n^2$ tasks
 - ▶ suppose that task (i, j, j) holds A_{ij} and task (i, j, i) holds B_{ij}
 - ▶ we will refer to tasks along i and j as task rows and columns
 - ▶ and tasks along k as *layers*
- ▶ the communication requirements among tasks are satisfied if we
 - ▶ broadcast the entries of the k^{th} column of A from task (i, j, j) to each task row in the k^{th} layer
 - ▶ broadcast the entries of the k^{th} row of B from task (i, j, i) to each task column of the k^{th} layer
 - ▶ form the result C_{ij} by the sum reduction of the values held by all the tasks layers k

Coarsening

- ▶ there are four natural ways to coarsen our $n \times n \times n$ fine grain tasks into p coarse grain tasks
 - ▶ by task rows: combine the $(n/p) \times n \times n$ tasks along a given task row
 - ▶ by task columns: combine the $n \times (n/p) \times n$ tasks along a given task column
 - ▶ partition the layers in a two dimensional grid by combining $(n/\sqrt{p}) \times (n/\sqrt{p}) \times n$ fine grain tasks
 - ▶ using three dimensional blocks by combining $(n/\sqrt[3]{p}) \times (n/\sqrt[3]{p}) \times (n/\sqrt[3]{p})$ tasks
- ▶ the two one dimensional coarsening strategies are similar
 - ▶ for row coarsening
 - ▶ each task needs only the part of A it already has, but needs all B entries
 - ▶ so global communication is required to broadcast the n^2/p entries of B held by each task
 - ▶ conversely, for column coarsening
 - ▶ each task needs only the parts of B that it already has, but it needs all of A
 - ▶ so global communication is required to broadcast the n^2/p entries of A held by each task
 - ▶ if accumulating A or B on each processor is not feasible, tasks can circulate portions of the array in a ring

Coarsening using a two dimensional grid

- ▶ block matrix multiplication has the same overall form as actual product, with scalar operations replaced by the matrix product of blocks!
- ▶ you should verify that

$$C_{ij} = \sum_{k=1}^{\sqrt{p}} A_{ik} B_{kj} \quad (5)$$

for $i, j = 1, \dots, \sqrt{p}$

- ▶ assume that task (i,j) has local access to block A_{ij} and B_{ij} and computes block C_{ij} of the result
- ▶ this requires all blocks A_{ik} and B_{kj} for $k = 1, \dots, \sqrt{p}$ to be communicated
 - ▶ first, a global broadcast of A blocks across each task row
 - ▶ followed by a global broadcast of B blocks across each task column
- ▶ memory requirements can be addressed by either of the following:
 - ▶ broadcast blocks of A across rows while circulating blocks of B across columns in lock step, so that they arrive at a given task at the same time
 - ▶ circulate blocks of A horizontally and blocks of B vertically, after an initial circular shift, so that blocks meet at a given task at the right time