

# ACM/CS 114

## Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Winter 2010

# LU factorization

- ▶ systems of linear equations are ubiquitous in numerical analysis
- ▶ let  $A$  be an  $n \times n$  matrix,  $b$  a known  $n$ -vector; we are looking for  $x$  such that

$$Ax = b \quad (1)$$

- ▶ a commonly used direct method for solving this system is to convert  $A$  into the product of a lower triangular matrix  $L$  with an upper triangular matrix  $U$

$$A = LU \quad (2)$$

known as  $LU$  factorization

- ▶ Eq. 1 becomes

$$LUx = b \quad (3)$$

which we can now solve in two simpler steps

$$Ly = b \quad (4)$$

$$Ux = y \quad (5)$$

where we first solve the lower triangular system by forward substitution, followed by solving the upper triangular system by back substitution to obtain  $x$

# LU by Gaussian elimination

- ▶ we can compute the  $LU$  factorization of  $A$  using Gaussian elimination

---

**Algorithm 1:** LU( $A$ )

---

```
1 for  $k = 1$  to  $n - 1$  do  
2   for  $i = k + 1$  to  $n$  do  
3      $L_{ik} = A_{ik} / A_{kk}$   
4   for  $j = k + 1$  to  $n$  do  
5     for  $i = k + 1$  to  $n$  do  
6        $A_{ij} = A_{ij} - L_{ik}A_{kj}$ 
```

---

which encodes  $L$  and  $U$  in place by overwriting  $A$

- ▶ Alg. 1 requires roughly  $n^3/3$  multiply-adds and  $n^2/2$  divisions
- ▶ we may also need *pivoting* to ensure numerical stability (and existence)
- ▶ Alg. 1 is one of many algorithms expressed essentially as a triply nested loop
  - ▶ the three indices can be ordered in any of  $3!$  ways, with totally different memory access patterns
  - ▶ in parallel, the  $kji$  and  $kji$  forms may be the most efficient

# Parallel $LU$ decomposition

- ▶ number fine grain tasks as  $(i,j)$  with  $i,j = 1, \dots, n$ ; each task
  - ▶ stores  $A_{ij}$
  - ▶ computes and stores  $U_{ij}$ , if  $i \leq j$
  - ▶ computes and stores  $L_{ij}$ , if  $i > j$

yielding a two dimensional array of  $n^2$  tasks

- ▶ no need to compute and store
  - ▶ the zeroes in the lower triangle of  $U$
  - ▶ the unit diagonal and the zeroes in the upper triangle of  $L$
- ▶ in order to create  $p$  coarse grain tasks we could combine
  - ▶  $n/p$  rows or columns of fine grain tasks
  - ▶  $(n/\sqrt{p}) \times (n/\sqrt{p})$  blocks of tasks

and map each one to a process

# Communication patterns for parallel $LU$ decomposition

---

**Algorithm 2:**  $LU(A, \text{task}=(i,j))$ 

---

```
1 for  $k = 1$  to  $\min(i,j) - 1$  do
2   recv  $A_{kj}$ 
3   recv  $L_{ik}$ 
4    $A_{ij} = A_{ij} - L_{ik}A_{kj}$ 
5 if  $i \leq j$  then
6   broadcast  $A_{ij}$  to  $(k,j), k = i + 1, \dots, n$ 
7 else
8   recv  $A_{jj}$ 
9    $L_{ij} = A_{ij}/A_{jj}$ 
10  broadcast  $L_{ij}$  to  $(i,k), k = i + 1, \dots, n$ 
```

---

# Row coarsening

- ▶ with one dimensional row coarsening
  - ▶ we forgo parallelism in updating rows
  - ▶ there is no need to broadcast the multipliers  $L_{ij}$  since each row is contained entirely within a task
  - ▶ we still need the vertical broadcasts of matrix rows to the tasks below

---

**Algorithm 3:** LU( $A$ , task= $(i,j)$ ) by rows

---

```
1 for  $k = 1$  to  $n - 1$  do
2   if  $k \in \text{myrows}$  then
3     broadcast  $\{A_{kj} : k \leq j \leq n\}$ 
4   else
5     recv  $\{A_{kj} : k \leq j \leq n\}$ 
6   for  $i \in \text{myrows}, i > k$  do
7      $L_{ik} = A_{ik} / A_{kk}$ 
8   for  $j = k + 1$  to  $n$  do
9     for  $i \in \text{myrows}, i > k$  do
10       $A_{ij} = A_{ij} - L_{ik}A_{kj}$ 
```

---

# Observations on row coarsening

- ▶ each task becomes idle as soon as it last row is completed
  - ▶ if rows are contiguous, a task may finish long before the overall computation is done
  - ▶ even worse, updating rows requires progressively less work with increasing row number
- ▶ we may improve concurrency and load balance
  - ▶ by assigning rows to tasks in a cyclic manner where row  $i$  is updated by task  $i \bmod p$
  - ▶ other mappings may be useful
- ▶ other improvements involve overlapping computation with communication
  - ▶ at step  $k$ , each task completes updating its portion of the remaining unreduced matrix before moving on to step  $k + 1$
  - ▶ however, the task that owns the  $k + 1$  row could broadcast it as soon as it becomes available, before moving on to the step  $k$  update
  - ▶ this *send ahead* strategy may grant other tasks earlier access to the data necessary to start working on the next step

# Column coarsening

---

**Algorithm 4:**  $\text{LU}(A, \text{task}=(i,j))$  by columns

---

```
1 for  $k = 1$  to  $n - 1$  do
2   if  $k \in \text{mycolumns}$  then
3     for  $i = k + 1$  to  $n$  do
4        $L_{ik} = A_{ik} / A_{kk}$ 
5       broadcast  $\{L_{ik} : k < i \leq n\}$ 
6   else
7     recv  $\{L_{ik} : k < i \leq n\}$ 
8   for  $i \in \text{mycolumns}, j > k$  do
9     for  $i = k + 1$  to  $n$  do
10       $A_{ij} = A_{ij} - L_{ik}A_{kj}$ 
```

---

- observations similar to row coarsening apply



## Block coarsening

---

**Algorithm 5:** LU( $A$ , task= $(i, j)$ ) by blocks

```

1 for  $k = 1$  to  $n - 1$  do
2   if  $k \in \text{myrows}$  then
3     broadcast  $\{A_{kj} : j \in \text{mycolumns}, j > k\}$  to all tasks in my task
      column
4   else
5     recv  $\{A_{kj} : j \in \text{mycolumns}, j > k\}$ 
6   if  $k \in \text{mycolumns}$  then
7     for  $i \in \text{myrows}, i > k$  do
8        $L_{ik} = A_{ik} / A_{kk}$ 
9     broadcast  $\{L_{ik} : i \in \text{myrows}, i > k\}$  to all tasks in my task row
10  else
11    recv  $\{L_{ik} : i \in \text{myrows}, i > k\}$ 
12  for  $j \in \text{mycolumns}, j > k$  do
13    for  $i \in \text{myrows}, i > k$  do
14       $A_{ij} = A_{ij} - L_{ik}A_{kj}$ 

```

## Observations on block coarsening

- ▶ each task becomes idle as soon as its last row and column are completed
  - ▶ if rows and columns are in contiguous blocks, a task may finish long before the overall computation is done
  - ▶ even worse, computing multipliers and updating blocks requires progressively less work with increasing row and column numbers
- ▶ we may improve concurrency and load balance
  - ▶ by assigning rows and columns to tasks in a cyclic manner where  $A_{ij}$  is assigned to task  $(i \bmod \sqrt{p}, j \bmod \sqrt{p})$
  - ▶ other mappings may be useful
- ▶ other improvements involve overlapping computation with communication
  - ▶ at step  $k$ , each task completes updating its portion of the remaining unreduced submatrix before moving on to step  $k + 1$
  - ▶ the broadcast of each segment of row  $k + 1$ , and the computation and broadcast of each segment of multipliers for step  $k + 1$ , can be initiated as soon as the relevant segments of row  $k + 1$  and column  $k + 1$  have been updated by their owners, before moving to completing the update for step  $k$
  - ▶ this *send ahead* strategy may grant other tasks earlier access to the data necessary to start working on the next step

# Pivoting

- ▶ the order of rows of  $A$  does not affect the solution to the system of equations
  - ▶ *partial pivoting* sorts the rows by the largest absolute value of the leading column of the remaining unreduced matrix
  - ▶ this choice ensures that the magnitude of the multipliers do not exceed 1, which
    - ▶ reduces amplification of round-off errors
    - ▶ ensures existence
    - ▶ improves numerical stability
- ▶ partial pivoting introduces a permutation matrix  $P$ , which leads to the factorization

$$PA = LU \tag{6}$$

which implies that the solution  $x$  is obtained through

$$Ly = Pb \tag{7}$$

$$Ux = y \tag{8}$$

with forward substitution in the lower triangular system, followed by back substitution in the upper triangular system

# Pivoting in parallel

- ▶ increased numerical stability costs increased parallel complexity and significant performance implications
- ▶ for one dimensional coarsening by column, the search for the pivot element requires no extra communication, but it is purely serial
  - ▶ once the pivot is found, the index of the pivot row must be communicated to the other tasks, and rows must be explicitly or implicitly interchanged in each task
- ▶ for coarsening by rows, the search for the pivot is parallel, but it requires communication among tasks and inhibits the overlapping of successive steps
  - ▶ if rows are explicitly interchanged, then only two tasks are involved
  - ▶ if rows are implicitly interchanged, changes to the assignment of rows to tasks are required, which has effects on concurrency and load balance
- ▶ in the presence of partial pivoting, column and row coarsening trade off on the relative speeds of computation versus communication
- ▶ with two dimensional coarsening, pivot search is parallel but requires communication among tasks along columns and destroys the possibility of overlapping successive steps

# Alternatives to pivoting

- ▶ various alternatives have been proposed
  - ▶ constraining pivoting to blocks of rows
  - ▶ pivoting when the multiplier exceeds a given threshold
  - ▶ pairwise pivoting
- ▶ these strategies are not foolproof, and trade off some stability and accuracy for speed

# Cholesky factorization

- ▶ when  $A$  is a positive definite symmetric matrix it has a Cholesky factorization

$$A = LL^T \quad (9)$$

with  $L$  a lower triangular matrix with positive entries along the diagonal

- ▶ so the linear system  $Ax = b$  can be solved through

$$Ly = b \quad (10)$$

$$L^T x = y \quad (11)$$

- ▶ the factorization is derived by equating corresponding entries of  $A$  with those of  $LL^T$  and generating them in the correct order
  - ▶ for example, in the  $2 \times 2$  case

$$\begin{bmatrix} A_{11} & A_{21} \\ A_{21} & A_{22} \end{bmatrix} = \begin{bmatrix} L_{11} & 0 \\ L_{21} & L_{22} \end{bmatrix} \begin{bmatrix} L_{11} & L_{21} \\ 0 & L_{22} \end{bmatrix} \quad (12)$$

yields

$$L_{11} = \sqrt{A_{11}} \quad L_{21} = A_{21}/L_{11} \quad L_{22} = \sqrt{A_{22} - L_{21}^2} \quad (13)$$

# Computing the Cholesky factorization

---

**Algorithm 6:** CHOLESKY(A)

---

```
1 for  $k = 1$  to  $n$  do
2    $A_{kk} = \sqrt{A_{kk}}$ 
3   for  $i = k + 1$  to  $n$  do
4      $A_{ik} = A_{ik}/A_{kk}$ 
5   for  $j = k + 1$  to  $n$  do
6     for  $i = j$  to  $n$  do
7        $A_{ij} = A_{ij} - A_{ik}A_{jk}$ 
```

---

- ▶ note that
  - ▶  $n$  square roots are required, all of positive numbers
  - ▶ only lower triangle of  $A$  is accessed, so the strict upper triangular part need not be stored
  - ▶  $A$  becomes  $L$  in place
  - ▶ the algorithm is stable so no pivoting is required
- ▶ it takes roughly half the number of  $LU$  operations: approximately  $n^3/6$  multiply-adds

# Parallelizing Cholesky

- ▶ number fine grain tasks as  $(i,j)$  with  $i,j = 1, \dots, n$ ; each task
  - ▶ stores  $A_{ij}$
  - ▶ computes and stores  $L_{ij}$ , if  $i \geq j$
  - ▶ computes and stores  $L_{ji}$ , if  $i < j$

yielding a two dimensional array of  $n^2$  tasks

- ▶ no need to compute and store the zero entries in the upper triangle



# Communication patterns for parallel Cholesky

---

**Algorithm 7:** CHOLESKY( $A$ , task= $(i, j)$ )

---

```
1 for  $k = 1$  to  $\min(i, j) - 1$  do
2   recv  $A_{kj}$ 
3   recv  $A_{ik}$ 
4    $A_{ij} = A_{ij} - A_{ik}A_{kj}$ 
5 if  $i = j$  then
6    $A_{ii} = \sqrt{A_{ii}}$ 
7   broadcast  $A_{il}$  to tasks  $(k, i)$  and  $(i, k)$ ,  $k = i + 1, \dots, n$ 
8 if  $i < j$  then
9   recv  $A_{ii}$ 
10   $A_{ij} = A_{ij}/A_{ii}$ 
11  broadcast  $A_{ij}$  to  $(k, j)$ ,  $k = i + 1, \dots, n$ 
12 if  $i > j$  then
13  recv  $A_{jj}$ 
14   $A_{ij} = A_{ij}/A_{jj}$ 
15  broadcast  $A_{ij}$  to  $(i, k)$ ,  $k = j + 1, \dots, n$ 
```

---

# Coarsening

- ▶ strategies very similar to  $LU$  factorization
  - ▶ one dimensional by row or column
  - ▶ two dimensional blocks

with column coarsening used most often in practice

- ▶ each choice of index in the outer loop yields different algorithm, named after the portion of the matrix that is updated by the basic operation in the inner loops
  - ▶ submatrix Cholesky: with  $k$  as the outer loop index, the inner loops perform a rank 1 update of the remaining unreduced submatrix, using the current column
  - ▶ column Cholesky: with  $j$  in the outer loop, inner loops compute the current column, using matrix-vector multiplies that accumulates the effects of previous columns
  - ▶ row Cholesky: with  $i$  in the outer loop, inner loops compute current row by solving a triangular system involving the previous rows

# Cholesky memory access patterns

