

ACM/CS 114

Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Spring 2012

Amortizing cost: containers

- ▶ the flexible solution is more expensive
 - ▶ this is sensible and, at some level, inevitable: flexibility implies runtime decisions
 - ▶ so we should be prepared to tolerate some slow down
- ▶ part of bad performance is caused by executing the *overhead* multiple times
- ▶ can we amortize the overhead over N , the total number of points?

Using a container for the sampling points

the base class, in `PointCloud.py`

```
1 class PointCloud(object):
2     """
3     The abstract base class for point generators
4     """
5
6     # interface
7     def points(self, n, box):
8         """
9         Generate {n} random points on the interior of {box}
10
11         parameters:
12             {n}: the number of points to generate
13             {box}: a pair of points on the plane that specify the major diagonal of
14                   rectangular region
15         """
16         raise NotImplementedError(
17             "class {.__name__!r} should implement 'points'".format(type(self)))
```

Filling a container with random points

in Mersenne.py

```
1 import random
2 from PointCloud import PointCloud
3
4 class Mersenne(PointCloud):
5     """
6     A point generator implemented using the Mersenne Twister random number
7     generator that is available as part of the python standard library
8     """
9
10    # interface
11    def points(self, n, box):
12        """
13        Generate {n} random points in the interior of {box}
14        """
15        # unfold the bounding box
16        intervals = tuple(zip(*box))
17        # create the container for the sample
18        sample = []
19        # loop over the sample size
20        while n > 0:
21            p = [ random.uniform(left, right) for left, right in intervals ]
22            sample.append(p)
23            n -= 1
24        # return the samples
25        return sample
```

Shape requires little modification

in Shape.py

```
1 class Shape(object):
2     """
3     The abstract base class for representations of geometrical regions
4     """
5
6     # interface
7     def interior(self, points):
8         """
9         Examine each point in {points} and return a container of only the
10        interior points
11        """
12        raise NotImplementedError(
```

The updated interior of a circle, in `Disk.py`

```
1 from Shape import Shape
2
3 class Disk(Shape):
4     """
5     A representation of a circular disk
6     """
7
8     # interface
9     def interior(self, points):
10        """
11        Build a container with members of {points} that fall within my disk
12        """
13        # precompute the frequently used values
14        r2 = self.radius**2
15        x0, y0 = self.center
16        # initialize the container of interior points
17        keep = []
18        # iterate over the given points and save the interior ones
19        for point in points:
20            x, y = point
21            dx = x - x0
22            dy = y - y0
23            # save the point if it is interior
24            if dx*dx + dy*dy <= r2:
25                keep.append(point)
26        # and return them to the caller
27        return keep
28
29    # meta methods
30    def __init__(self, radius=1.0, center=(0.0, 0.0)):
31        self.radius = radius
32        self.center = center
33        return
```

The driver for the container based solution

```
1
2 def gauss():
3     """
4     The driver for the container based implementation
5     """
6     from Disk import Disk
7     from Mersenne import Mersenne
8
9     # inputs
10    N = 10**5
11    box = [(0,0), (1,1)]
12    # the point cloud generator
13    generator = Mersenne()
14    # the region of integration
15    disk = Disk(center=(0,0), radius=1)
16
17    # the integration algorithm
18    # build the point sample
19    sample = generator.points(N, box)
20    # count the interior points
21    interior = len(disk.interior(sample))
22
23    # print out the estimate of  $\pi$ 
24    print("pi: {0:.8f}".format(4*interior/N))
25    return
```

The performance of the container solution

N	C++ $t(\text{sec})$	python $t(\text{sec})$	naïve OO $t(\text{sec})$	containers $t(\text{sec})$
10^0	.002	.014	.014	.014
10^1	.002	.014	.014	.014
10^2	.002	.014	.014	.015
10^3	.002	.015	.020	.019
10^4	.004	.027	.078	.063
10^5	.026	.144	.625	.504
10^6	.230	1.265	6.242	5.925
10^7	2.277	12.624	61.583	188.318
10^8	22.749	130.430		
10^9	227.735			

Generators

- ▶ the container solution has a certain elegance
 - ▶ the creation of the container and its use are separated
 - ▶ we pay less overhead per point
 - ▶ with our points in containers, we can use the fast *functional* routines to operate on them
- ▶ but, we have now uncovered a new source of cost: managing memory
 - ▶ it appears that building large lists is expensive
 - ▶ can we avoid building the container all together?
- ▶ generators!

Generating random points

in Mersenne.py

```
1 import random
2 from PointCloud import PointCloud
3
4 class Mersenne(PointCloud):
5     """
6     A point generator implemented using the Mersenne Twister random number genera
7     available as part of the python standard library
8     """
9
10    # interface
11    def points(self, n, box):
12        """
13        Generate {n} random points in the interior of {box}
14        """
15        # unfold the bounding box
16        intervals = tuple(zip(*box))
17        # loop over the sample size
18        while n > 0:
19            p = [ random.uniform(*interval) for interval in intervals ]
20            yield p
21            n -= 1
22
23    return
```

The modified Disk.py

```
1 from Shape import Shape
2
3 class Disk(Shape):
4     """
5     A representation of a circular disk
6     """
7
8     # interface
9     def interior(self, points):
10        """
11        Predicate that filters out points that are not in my interior
12        """
13        # precompute the frequently used values
14        r2 = self.radius**2
15        x0, y0 = self.center
16        # iterate over the given points and return the interior ones
17        for point in points:
18            x, y = point
19            dx = x - x0
20            dy = y - y0
21            if dx*dx + dy*dy <= r2:
22                yield point
23        # all done
24        return
25
26
27 # meta methods
28 def __init__(self, radius=1.0, center=(0.0, 0.0)):
29     self.radius = radius
30     self.center = center
31     return
```

The driver for the generator based solution

```
1
2 def gauss():
3     """
4     The driver for the generator based implementation
5     """
6     from Disk import Disk
7     from Mersenne import Mersenne
8
9     # inputs
10    N = 10**7
11    box = [(0,0), (1,1)]
12    # the point cloud generator
13    generator = Mersenne()
14    # the region of integration
15    disk = Disk(center=(0,0), radius=1)
16
17    # the integration algorithm
18    # build the point sample
19    sample = generator.points(N, box)
20    # count the interior points
21    interior = count(disk.interior(sample))
22
23    # print out the estimate of  $\pi$ 
24    print("pi: {0:.8f}".format(4*interior/N))
25    return
26
27
28 def count(iterable):
29     """
30     Count the entries of iterable
31     """
32     counter = 0
33     for item in iterable:
34         counter += 1
35     return counter
```

The performance of the generator solution

N	C++ $t(\text{sec})$	python $t(\text{sec})$	naïve OO $t(\text{sec})$	containers $t(\text{sec})$	generators $t(\text{sec})$
10^0	.002	.014	.014	.014	.014
10^1	.002	.014	.014	.014	.014
10^2	.002	.014	.014	.015	.014
10^3	.002	.015	.020	.019	.018
10^4	.004	.027	.078	.063	.053
10^5	.026	.144	.625	.504	.401
10^6	.230	1.265	6.242	5.925	3.780
10^7	2.277	12.624	61.583	188.318	38.242
10^8	22.749	130.430			
10^9	227.735				

Representing integrands

in `Functor.py`

```
1 class Functor(object):
2     """
3     The abstract base class for function objects
4     """
5
6     # interface
7     def eval(self, points):
8         """
9         Evaluate the function at the supplied points
10        """
11        raise NotImplementedError(
12            "class {.__name__!r} should implement 'eval'".format(type(self)))
```

Constant functions

in `Constant.py`

```
1 from Functor import Functor
2
3 class Constant(Functor):
4     """
5     A representation of constant functions
6     """
7
8     # interface
9     def eval(self, points):
10         """
11         Compute the value of the function
12         """
13         # cache the constant
14         constant = self.constant
15         # return the constant regardless of the evaluation point
16         for point in points: yield constant
17         # all done
18         return
19
20     # meta methods
21     def __init__(self, constant):
22         self.constant = constant
23         return
```

A non-trivial integrand

```
1 from Functor import Functor
2
3 class Gaussian(Functor):
4     """
5     An implementation of the normal distribution with mean  $\mu$  and variance  $\sigma^2$ 
6     """
7
8     # interface
9     def eval(self, points):
10        """
11        Compute the value of the gaussian
12        """
13        # access the math symbols
14        from math import exp, sqrt, pi
15        # cache the shape information
16        mean = self.mean
17        spread = self.spread
18        # precompute the normalization factor and the exponent scaling
19        normalization = 1 / sqrt(2*pi) / spread
20        scaling = 2 * spread**2
21        # loop over points and yield the computed value
22        for p in points:
23            # compute the norm |p - mean|^2
24            # this works as long as {p} and {mean} have the same length
25            r2 = sum((p_i - mean_i)**2 for p_i, mean_i in zip(p, mean))
26            # yield the value at the current p
27            yield normalization * exp(- r2/scaling)
28        # all done
29        return
30
31    # meta methods
32    def __init__(self, mean, spread):
33        self.mean = mean
34        self.spread = spread
35        return
```


The Monte Carlo integrator

```
1 import operator, functools
2
3 def gauss():
4     """
5     The driver for the generator based implementation
6     """
7     from Disk import Disk
8     from Gaussian import Gaussian
9     from Mersenne import Mersenne
10
11     # inputs
12     N = 10**7
13     box = [(-1,-1), (1,1)]
14     B = functools.reduce(operator.mul, ((right-left) for left,right in zip(*box)))
15     # the point cloud generator
16     generator = Mersenne()
17     # the region of integration
18     disk = Disk(center=(0,0), radius=1)
19     # the integrand
20     gaussian = Gaussian(mean=(0,0), spread=1/3)
21
22     # the integration algorithm
23     # build the point sample
24     sample = generator.points(N, box)
25     # select the interior points
26     interior = disk.interior(sample)
27     # compute the integral
28     integral = B/N * sum(gaussian.eval(interior))
29
30     # print out the estimate of the integral
31     print("integral: {0:.8f}".format(integral))
32     return
```