

# ACM/CS 114

## Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Spring 2012

# Recap: what we know so far

- ▶ `pyre` components are evolved python objects
  - ▶ the factories have family names, the instances have names
  - ▶ these names are unique strings in hierarchical namespaces delimited by periods
  - ▶ collections of components form packages *implicitly*, based on the topmost level in their namespace
- ▶ components have properties that are under the control of the *user*
  - ▶ they look and behave like regular attributes
  - ▶ they are *typed* to enable conversions from strings
  - ▶ they have default values and other metadata
- ▶ configuration is partly about assigning values to component properties
  - ▶ a requirement for supporting user interfaces
  - ▶ intuitive syntax for the command line
  - ▶ simple configuration files inspired by the Microsoft Windows `.ini` format
- ▶ configuration is automatically handled by the framework and requires no explicit involvement on the part of the component author

# A simple component

```
1 import pyre
2
3 class Disk(pyre.component, family="gauss.shapes.disk"):
4
5     # public state
6     radius = pyre.properties.float()
7     radius.default = 1
8     radius.doc = 'the radius of the disk'
9
10    center = pyre.properties.array()
11    center.default = (0,0)
12    center.doc = 'the location of the center of the circle'
13
14    # interface
15    @pyre.export
16    def interior(self, points):
17        """
18        Filter {points} by removing exterior ones
19        """
20        ...
```

# A simple configuration file

- ▶ a file named `gauss.cfg` with the following

```
1  [ gauss.shapes.disk ] ; the family name
2  radius = 2
3  center = (-1,-1)
4
5  [ disk1 ] ; the name of an instance
6  center = (-1,1) ; leave {radius} alone
7
8  [ disk2 ] ; the name of another instance
9  radius = .5
10 center = (1,1)
```

could get loaded automatically when `Disk` is first imported, or loaded by naming it explicitly on the command line

- ▶ the `.cfg` extension allows the framework to deduce which configuration file parser should be used to process the contents

# Configuration files

- ▶ currently, there are two file formats for configuration information
  - ▶ `.cfg`: the format in the examples
  - ▶ `.pml`: an XML based format that is a bit more powerful but not as user friendly
- ▶ `pyre` looks for configuration files in the following places
  - ▶ explicitly provided on the command line

```
1 gauss.py --config=sample.cfg
```

- ▶ the current directory
- ▶ the `.pyre` subdirectory of the current user's home directory
- ▶ a special subdirectory wherever `pyre` is installed

in order of priority

- ▶ settings on the command line have the highest priority, and override each other from left to right
- ▶ when a property is assigned a value multiple times, the highest priority setting wins
  - ▶ the framework keeps track of all changes in the value of properties and the source of the assignment, so if a property doesn't end up with the value you expected, you can get its complete history

# Properties

- ▶ properties make sense for both classes and instances
  - ▶ the class holds the default value that gets used in case the component instance does not have explicit configuration
  - ▶ each instance gets its own private value when it gets configured
  - ▶ identical to regular python attributes
- ▶ there is support for
  - ▶ simple types: `bool`, `int`, `float`, `str`
  - ▶ containers: `tuple`, `array`
  - ▶ higher level: `date`, `time`, `inputfile`, `outputfile`, `inet`
  - ▶ units: `dimensional`
  - ▶ easy enough to implement your own; the requirements are very simple
- ▶ metadata:
  - ▶ doc: a simple and short documentation string
  - ▶ default: the default value, in case the user doesn't supply one
  - ▶ converters: a chain of preprocessors of the string representation
  - ▶ normalizers: a chain of post-processors of the converted value
  - ▶ validators: a tuple of predicates that get called to ensure the property value satisfies the specified constraints
  - ▶ you can add your own; the framework passes them through to your component

- ▶ dimensional properties have units
- ▶ the low level support is in `pyre.units`
  - ▶ full support for all SI base and derived units
  - ▶ all common abbreviations and names from alternative systems of units
  - ▶ correct arithmetic; proper handling of functions from `math`

```
1 from math import cos
2 from pyre.units.SI import meter, second, radian
3
4 A = 2.5 * meter
5 t = 1.5 * second
6  $\omega$  = 4.2 * radian/second
7
8 x = A * cos( $\omega$  * t)
```

if the units in the argument to `cos` do not cancel, leaving a pure float behind, an exception is raised; `x` has dimensions of meters

# Connecting components

the real power is in wiring components to other components

```
1 import pyre
2
3 class MonteCarlo(pyre.component, family="gauss.integrators.montecarlo"):
4     """
5     A Monte Carlo integrator
6     """
7
8     # public state
9     samples = pyre.property.int(default=10**5)
10    region = ???
11    ...
```

- ▶ MonteCarlo should be able to specify what constitutes an acceptable region
- ▶ Disk should be able to advertise itself as being an acceptable region
- ▶ the user should have natural means for specifying that she wants to wire an instance of Disk as the region of integration
- ▶ and be able to configure that particular instance of Disk in a natural manner
- ▶ the framework should check the consistency of this assignment



# Interfaces

the component version of our abstract base class is the *interface*

```
1 import pyre
2
3 class Shape(pyre.interface, family="gauss.shapes"):
4     """
5     The obligations of implementations of geometrical shapes
6     """
7
8     # my default implementation
9     @classmethod
10    def default(cls):
11        """
12        The default {Shape} implementation
13        """
14        # use a box
15        from .Box import Box
16        return Box # if you return an instance, it will be shared by all...
17
18    # interface
19    @pyre.provides
20    def measure(self):
21        """
22        Compute my measure: length, area, volume, etc
23        """
24
25    @pyre.provides
26    def interior(self, points):
27        """
28        Filter out {points} that are on my exterior
29        """
```

pyre prohibits the instantiation of interfaces, so you don't have to worry about the implementation of the methods

# Declaring compatibility with an interface

Disk can inform the framework that it intends to implement Shape

```
1 import pyre
2 from Shape import Shape
3
4 class Disk(pyre.component, family="gauss.shapes.disk", implements=Shape):
5     """
6     A representation of a circular disk
7     """
8
9     ...
```

an exception is raised if `Disk` does not conform fully to `Shape`

- ▶ missing methods or missing attributes

also, proper namespace design simplifies many things for the user

- ▶ `Shape` declared its family as `gauss.shapes`
- ▶ `Disk` declared its family as `gauss.shapes.disk`

we'll see how later when it's time to put all this together

# Specifying assignment requirements

MonteCarlo can now specify it expects a Shape compatible object to be assigned as its region of integration

```
1 import pyre
2 from Shape import Shape
3
4 class MonteCarlo(pyre.component, family="gauss.integrators.montecarlo"):
5     """
6     A Monte Carlo integrator
7     """
8
9     # public state
10    samples = pyre.property.int(default=10**5)
11    region = pyre.facility(interface=Shape)
12    ...
```

the default value is whatever Shape returns from its default class method

# Component specification

- ▶ and now for the real trick: converting some string provided by the user into a live instance of `Disk`, configuring it, and attaching it to some `MonteCarlo` instance
- ▶ the syntax is motivated by URI, the universal resource identifiers of the web; the general form is

```
1 <scheme>://<authority>/<path>#<identifier>
```

where most of the segments are optional

- ▶ if your component is accessible from your python path, you could specify

```
1 import:gauss.shapes.disk
```

- ▶ if your component instance is somewhere on your disk, you would specify

```
1 file:/tmp/shapes.py/disk
```

- ▶ in either case, `disk` is expected to be a name that resolves into a component class, a component instance or be a callable that returns one of these

# The user does the wiring

with our definition of `MonteCarlo`, an appropriately structured package `gauss` on the python path, and the following configuration file

```
1 [ gauss.shapes.disk ] ; change the default values for all disks
2 radius = 1
3 center = (0,0)
4
5 [ mc ] ; configure our Monte Carlo integrator instance
6 region = import:gauss.shapes.disk
7 ...
```

the following python code in some script `sample.py`

```
1 ...
2 mc = MonteCarlo(name="mc")
3 ...
```

builds a `MonteCarlo` instance and configures it so that its region of integration is a `Disk` instance; similarly, from the command line

```
1 sample.py --mc.region=import:gauss.shapes.disk
```

or, thanks to the consistency in our namespace layout, simply

```
1 sample.py --mc.region=disk
```