

ACM/CS 114

Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Spring 2012

Extensions

- ▶ one of the major strengths of python is the ease with which you can get access to code written in low level languages

- ▶ for leverage, reuse, performance

the interpreter itself is written in C, so you can get access to code written in any language that is link compatible

- ▶ python provides
 - ▶ a C library that grants you access to nearly every aspect of the interpreter
 - ▶ a simple set of rules for constructing shared objects that can be imported, just like any other python module
- ▶ typically, the solution to this problem has three layers
 - ▶ the archive and headers of the library you want to expose to python
 - ▶ the interface layer that translates data from python to C, which becomes the importable shared module
 - ▶ a layer in pure python that provides an object oriented veneer to the library; in our case this layer will also double as the component specification
- ▶ if tempted to collapse some of these, *resist*

The bottom layer

we would like use the random number generators in GSL; to get a feeling about what is involved, recall our earlier C++ solution

```
9 #include <iostream>
10 #include <gsl/gsl_rng.h>
11
12 int main(int, char*[]) {
13     // the sample size
14     const int N = 1.0e7;
15     // initialize the counters
16     int interior = 0;
17     // allocate a random number generator
18     gsl_rng * generator = gsl_rng_alloc(gsl_rng_ranlxs2);
19     // integrate by sampling some number of times
20     for (int i=0; i<N; ++i) {
21         // create a random point
22         double x = gsl_rng_uniform(generator);
23         double y = gsl_rng_uniform(generator);
24         // check whether it is inside the unit quarter circle
25         if ((x*x + y*y) <= 1.0) { // no square roots
26             // update the interior point counter
27             interior++;
28         }
29     }
30     // print the result
31     std::cout << "pi: " << 4. * interior / N << std::endl;
32     // free the generator
33     gsl_rng_free(generator);
34     // all done
35     return 0;
36 }
```

we must allocate, use and free a `gsl_rng`

The top layer

for plug-and-play, we must build a component; perhaps something similar to Mersenne:

```
1 import pyre, itertools, gsl
2 from .PointCloud import PointCloud
3
4 class GSL(pyre.component, family="gauss.meshes.gsl", implements=PointCloud):
5     """
6     A point generator implemented using the large set of random number
7     generators available as part of the gnu scientific library (GSL)
8     """
9
10    # interface
11    @pyre.export
12    def points(self, n, box):
13        """
14        Generate {n} random points in the interior of {box}
15        """
16        # unfold the bounding box
17        intervals = tuple(box.intervals()) # realize, so we can reuse in the loop
18        # loop over the sample size
19        while n > 0:
20            # make a point and yield it
21            yield tuple(itertools.starmap(self.rng.uniform, intervals))
22            # update the counter
23            n -= 1
24        # all done
25        return
```

where `self.rng` is a *handle* for the `gsl_rng` pointer

Adding the new component to our package

in `gauss/meshes/__init__.py`

```
9 """
10 Package that contains the implementations of point clouds
11 """
12
13 # the interfaces
14 from .PointCloud import PointCloud as cloud
15
16 # the components
17 from .GSL import GSL as gsl
18 from .Mersenne import Mersenne as mersenne
```

the only reason to tuck this in `gauss.meshes` is the convenience of using the shorter name when configuring

```
1 [ mc ] ; configure our Monte Carlo integrator instance
2 samples = 10**6
3 mesh = gsl
4 region = ball
5 integrand = constant
```

access to GSL could have been provided by a user of the `gauss` package just as easily, without any access to our source code

Module access

- ▶ python extensions live in dynamically loaded libraries
 - ▶ DLLs on windows, shared objects on unix
 - ▶ building one is platform dependent; compilers have the right command line flags
- ▶ making an extension requires a *visible entry point* in a shared object, both of whose names python can infer correctly so that the statement

```
1 import gsl
```

can be converted into a search for a particular file, followed by a lookup of a particular symbol

- ▶ e.g. on unix, acceptable names for the `gsl` extension module are `gsl.so` or `gslmodule.so`
- ▶ looking for extensions is just another step in the sequence applied to the folders on `sys.path`, so the same rules apply as looking for regular modules

The visible entry point

and here is the definition of the entry point

```
1 // initialization function for the gsl module
2 // *must* be called PyInit_gsl
3 PyMODINIT_FUNC PyInit_gsl()
4 {
5     // create the module
6     PyObject * module = PyModule_Create(&gsl::module_definition);
7     // check whether module creation succeeded
8     if (!module) {
9         // otherwise, raise an ImportError
10        return 0;
11    }
12    // we have an initialized module
13    // set the error handler
14    gsl_set_error_handler(&errorHandler);
15    // initialize the table of known random number generators
16    gsl::rng::initialize();
17
18    // return the newly created module
19    return module;
20 }
```

python will call this function to create the module object, if the dynamic loading of the extension is successful

Building a table of known generators

```
1 namespace gsl {
2     namespace rng {
3         typedef std::map<std::string, const gsl_rng_type *> map_t;
4         static map_t generators;
5     }
6 }
7
8 // initialization of the known generators
9 void
10 gsl::rng::initialize()
11 {
12     // iterate over the registered names
13     const gsl_rng_type **current = gsl_rng_types_setup();
14     while (*current != 0) {
15         // add each one to my map
16         gsl::rng::generators[(*current)->name] = *current;
17         // get the next one
18         current++;
19     }
20
21     return;
22 }
```


Module definition

the identifier `module_definition` refers to a struct

```
1 // the module documentation string
2 const char * const __doc__ = "sample module documentation string";
3 // the module definition structure
4 PyModuleDef gsl::module_definition = {
5     // header
6     PyModuleDef_HEAD_INIT,
7     // the name of the module
8     "gsl",
9     // the module documentation string
10    __doc__,
11    // size of the per-interpreter state of the module
12    -1, // don't touch
13    // the methods defined in this module
14    gsl::module_methods
15 };
```