

ACM/CS 114

Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Winter 2012

Algorithms

- ▶ Informally, an algorithm can be viewed as
 - ▶ a well-defined computational procedure that
 - ▶ takes a set of values as input
 - ▶ produces a set of values as output
 - ▶ a solution to a computational problem
 - ▶ whose statement specifies the intended relationship between inputs and outputs
 - ▶ and the algorithm being the specific computational procedure that achieves this relationship
- ▶ the prototypical computational problem is *sorting*
 - ▶ problem specification
 - ▶ input: a sequence S of n numbers (s_0, s_1, \dots, s_n)
 - ▶ output: a permutation S' of the input sequence $(s'_0, s'_1, \dots, s'_n)$
 - ▶ constraint: the elements of the output sequence must satisfy
$$s'_0 \leq s'_1 \leq \dots \leq s'_n$$
 - ▶ problem instance: $S = (0, \pi, 1, e, 2, 16)$
 - ▶ invalid input: $(0, i, 1)$
 - ▶ why is this bad? which implicit property of S does it violate? what is the set of valid inputs?

Correctness

- ▶ once again informally, an algorithm is *correct* if
 - ▶ it terminates for all valid input
 - ▶ upon termination on valid input, the output satisfies the constraints expressed in the problem statement
- ▶ equivalently, we say that the algorithm *solves* the computational problem
- ▶ after correctness has been established, algorithms are classified according to their demands on computational resources
 - ▶ running time complexity
 - ▶ a measure of the number of instructions necessary to solve the problem
 - ▶ and, occasionally
 - ▶ amount of auxiliary storage
 - ▶ network bandwidth or other communication infrastructure requirements
 - ▶ for parallel algorithms: speedup and efficiency
- ▶ algorithms are often specified using *pseudocode*
 - ▶ a loose language with mostly notational constraints
 - ▶ a mixture of reasonable looking code with whatever expressive method makes the point clear
 - ▶ hence, the use of human languages to convey meaning that might be too difficult to code up, or would obscure the point, is perfectly acceptable

A sorting algorithm

Algorithm 1: INSERTION-SORT(S)

```
1 for  $j \leftarrow 2$  to  $\text{length}[S]$  do
2    $\text{key} \leftarrow S[j]$ 
3    $i \leftarrow j - 1$ 
4   while  $i > 0$  and  $S[i] > \text{key}$  do
5      $S[i + 1] \leftarrow S[i]$ 
6      $i \leftarrow i - 1$ 
7    $S[i + 1] \leftarrow \text{key}$ 
```

- ▶ valid inputs:
 - ▶ empty sequence, singlet, other sequences of finite length
 - ▶ what kinds of objects in S ?
- ▶ walk through it by hand with $S = (5, 2, 4, 6, 1, 3)$

Pseudocode conventions

- ▶ the symbol “▷” indicates a comment through to the end of the line
- ▶ block structure is indicated by the indentation level
- ▶ all variables are local; no global variables, unless explicitly marked
- ▶ $i \leftarrow j \leftarrow k$ assigns the rightmost expression to all the other variables
- ▶ indexing: $S[i]$; slicing: $S[i..j]$
- ▶ conditionals, looping constructs, function calls should be familiar
- ▶ compound objects have attributes or fields that are referenced using indexing, e.g. $length[S]$
- ▶ variables assigned to objects or containers are references
- ▶ parameters passed to procedures *by assignment*

Python implementation

- ▶ direct translation of pseudocode in python, with no attempt to improve

```
1 def insertion_sort(S):  
2     for j in range(1, len(S)):  
3         key = S[j]  
4         i = j-1  
5         while i>=0 and S[i]>key:  
6             S[i+1] = S[i]  
7             i = i-1  
8         S[i+1] = key
```

- ▶ minor adjustments to loop indices are required since python lists are zero based

Analyzing algorithms

- ▶ algorithm analysis is the computation of resource requirements
 - ▶ memory, communication bandwidth, *computational time*
- ▶ need a model for the implementation environment
 - ▶ RAM: *random access machine*
 - ▶ an abstraction of a single processor sequential execution machine that has access to a single block of memory with uniform access cost
 - ▶ even though the model is extremely simple, algorithm analysis remains a hard problem, full of subtleties
- ▶ in general, we seek to relate the running time to input size
 - ▶ definition of input size is problem dependent – could be number of items to sort, or number of grid points in a mesh, etc.
 - ▶ running time is proportional to the number of primitive steps executed
 - ▶ different lines have different costs
 - ▶ but each execution of a given line is assumed to cost the same
- ▶ *exercise*: decorate Alg. 1 with the number of times each line is executed

Runtime complexity of INSERTION-SORT

- ▶ summing up the number of times each line of Alg. 1 gets executed:

$$T(n) = c_2n^2 + c_1n + c_0$$

where the c_i are constants related to the cost of the various lines


- ▶ note that it is a quadratic function of n
- ▶ best case: input S is already sorted
- ▶ worst case: input S is reverse-sorted
- ▶ average case: assume a *random* S and compute an expectation value for the number of executions of each line
 - ▶ still a quadratic function of n
- ▶ we quantify the run time complexity of INSERTION-SORT by saying that it is asymptotically bound by n^2
 - ▶ concentrating on the highest power of n
 - ▶ disregarding the multiplicative constants that are strongly dependent on the execution model, rather than the quality of the algorithm

Asymptotic bounds

- ▶ most often, run time complexity analysis is reduced to constructing asymptotic bounds on execution time as the input size $n \rightarrow \infty$
 - ▶ i.e., finding a simpler function of the input size with similar behavior for large n
- ▶ we say that $f = \Theta(g)$ if there are constants c_1, c_2 such that

$$c_1 g(n) \leq f(n) \leq c_2 g(n)$$

for sufficiently large n



Designing algorithms

- ▶ INSERTION-SORT is *incremental*:
 - ▶ having sorted $S[i..j]$, put $S[j]$ in its proper place
 - ▶ how would you break this up into tasks that can be executed in parallel?
- ▶ one alternative is *divide-and-conquer*: MERGE-SORT
 - ▶ *divide*: split S into two parts of roughly equal length
 - ▶ *conquer*: sort the subsequences recursively
 - ▶ *combine*: merge the two sorted subsequences to produce the sorted output

Sorting by divide-and-conquer

Algorithm 2: MERGE-SORT(S, p, r)

```
1 if  $p < r$  then  
2    $q \leftarrow \lfloor (p + r)/2 \rfloor$   
3   MERGE-SORT( $S, p, q$ )  
4   MERGE-SORT( $S, q + 1, r$ )  
5   MERGE( $S, p, q, r$ )
```

- ▶ *exercise*: write MERGE; can be done in $\Theta(r - p + 1)$
- ▶ analysis of running time:
 - ▶ involves solving a recurrence relation
 - ▶ worst case:

$$T(n) = \left\{ \begin{array}{ll} \Theta(1) & \text{if } n = 1 \\ 2T(n/2) + \Theta(n) & \text{if } n > 1 \end{array} \right\} \rightarrow \Theta(n \log n)$$

- ▶ is this a better candidate for parallel sorting?