

# ACM/CS 114

## Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Winter 2012

# Sources of complexity

- ▶ problem size
  - ▶ cpu time, amount of memory, cpu count
- ▶ project size
  - ▶ asset complexity: lines of code, entry points, files
  - ▶ dependencies: number of modules, third party libraries
  - ▶ run-time complexity: number of object types and instances
- ▶ project longevity
  - ▶ duty cycle, life cycle
  - ▶ cost/benefit of reuse
  - ▶ managing change: people, software, hardware, technology
- ▶ locality of needed resources
  - ▶ computation and persistence: where, how, when, who
- ▶ usability
  - ▶ learning curve, ease of access, interfaces, security, etc.

## Risk mitigation

- ▶ verification and validation
  - ▶ invariants, consistency checks, code introspection
  - ▶ benchmark problems
  - ▶ reality checks
- ▶ key software practices
  - ▶ source control: e.g. cvs; svn; git, hg, bzt
  - ▶ uniform build procedure: platforms, compilers, dependencies, but also build types
  - ▶ testing: unit, component, application; mutation, regression; *automate*
  - ▶ tracking of bugs and feature requests: gnats, bugzilla, trac
  - ▶ documentation: doxygen, epydoc; wiki
  - ▶ code reviews
- ▶ project coherence
  - ▶ design methodology
  - ▶ flexible software architecture
  - ▶ coding and documentation standards
  - ▶ regular meetings

# Remote access to computational resources

- ▶ `ssh`, the *secure shell*:
  - ▶ currently the most flexible and secure method
  - ▶ authentication, strong data encryption
  - ▶ interactive sessions, remote shell services, data transfer, command execution
- ▶ authentication requires a *private* and a *public* key
  - ▶ the private key is protected by a *passphrase*; must be long enough to be difficult to guess, but memorable because it is not recoverable
  - ▶ you only need one private key, and it only needs to be installed on machines whose keyboard you actually touch
  - ▶ you install your public key on every machine to which you require access
    - ▶ on unix machines: in `~/.ssh/authorized_keys`
  - ▶ keys are created using `ssh-keygen`; on windows there is PuTTY
- ▶ most installations support `ssh-agent`
  - ▶ a wrapper around your interactive session that asks you for your passphrase, loads your private key in memory, and negotiates authentication with `ssh` servers automatically

# Key generation with PuTTY

- ▶ the PuTTY screen before and after a key has been generated



# Bazaar – setup

- ▶ we are going to use `bzr`, a source control system, for both homework and final project submissions
- ▶ installation instructions, documentation and tutorials available from <http://bazaar.canonical.com>
- ▶ bazaar keeps track of authors in its revision history, so it's a good idea to give it your name and an email address

```
1 #> bzr whoami "Michael Aivazis <aivazis@caltech.edu>"
```

- ▶ you should verify that it worked by asking

```
1 #> bzr whoami
2 Michael Aivazis <aivazis@caltech.edu>
```

- ▶ the discussion here is specialized to the class repository, but the ideas generalize easily



# Getting information from the class server

- ▶ `bzr` is a *distributed* source control system: each repository has a copy of the entire change history, and can evolve independently
- ▶ you are responsible for sharing information among repositories by invoking the repository synchronization commands *explicitly* and providing the location of the remote repository
- ▶ populate your new `acm114` repository by *pulling* from the class server

```
1 #> bzr pull --remember bzr+ssh://acm114@acm114.caltech.edu
2 +N assignment-1/
3 +N assignment-1/20120120.pdf
4 +N assignment-1/20120120.tex
5 +N assignment-1/setup.tex
6 All changes applied successfully.
7 Now on revision 1.
```

- ▶ `--remember` asks `bzr` to use the class repository as the default: afterwards, you can just `bzr pull` without mentioning the source
- ▶ `bzr+ssh` is the connection scheme that employs `ssh` for authentication
- ▶ `acm114@acm114.caltech.edu` is the address of your homework repository on the class server



# Working in the repository

- ▶ one of the problems in the first assignment asks you to write a routine that computes the first few terms of some series in a language of your choice; let's assume that you picked `python`, and you created a file `dilog.py`

```
1 #> cd assignment-1
2 #> vi dilog.py
```

- ▶ you can inform `bzr` that you would like to place `dilog.py` under source control

```
1 #> bzr add dilog.py
2 adding assignment-1/dilog.py
```

- ▶ let's check what `bzr` knows about the repository

```
1 #> bzr status
2 added:
3   assignment-1/dilog.py
```

## Committing and submitting

- after *exhaustive* editing and testing, you are exhausted; time to take a break:

```
1 #> bzr commit -m "assignment-1: first draft of the dilog routine"
2 #> dilog.py
3 Committing to: /home/mga/courses/acml14
4 added assignment-1/dilog.py
5 Committed revision 2.
```

- ▶ rinse, repeat...
- ▶ time to *push* your work to the course server

```
1 #> bzr push --remember bzr+ssh://acm114@acm114.caltech.edu
2 ...
3 Pushed up to revision 7
```

- ▶ again, the `--remember` stores the location of the remote repository so that subsequent push operations do not need to specify it