

ACM/CS 114

Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Winter 2012

Motivations for going parallel

- ▶ why bother?
 - ▶ *speed*: there are fundamental limits to the processing power of a single processor
 - ▶ *throughput*: time to solution is critical for many problems
 - ▶ *size*: high resolution requires lots of memory
 - ▶ *availability*: the tool exists, use it
- ▶ but be careful
 - ▶ the commercial market is unstable
 - ▶ the computing environment is somewhat primitive
 - ▶ software packages and libraries are emerging slowly
 - ▶ parallel programming is not hard, but it requires *discipline*

Taxonomy

- ▶ an early classification of computer systems focused on the relation between *instruction streams* and *data streams*:
 - ▶ SISD: single instruction, single data
 - ▶ SIMD: single instruction, multiple data
 - ▶ MIMD: multiple instruction, multiple data
- ▶ SISD describes conventional serial computers
 - ▶ the programming model only: the hardware has moved on...
- ▶ SIMD and MIMD are the traditional models for parallel machines
- ▶ MIMD systems are often programmed in SPMD mode: single *program*, multiple data
 - ▶ when the parallel environment provides a *naming scheme* for the instruction streams, such as processor id, or task name

Architectural issues

- ▶ *control*: SIMD vs. MIMD
- ▶ *coördination*: synchronous vs. asynchronous
- ▶ *memory organization*: private vs. shared
- ▶ *address space*: local vs. global
- ▶ *memory access*: uniform vs. non-uniform
- ▶ *granularity*: the power of each processor
- ▶ *scalability*: dependence on the number of processors
- ▶ *interconnect*: topology, routing, switching

Tradeoffs

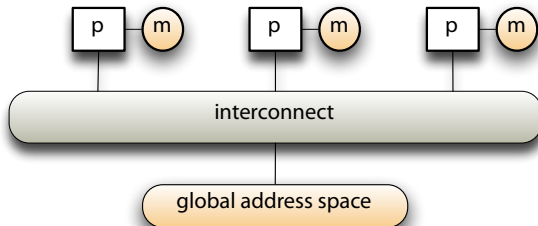
	shared memory	distributed memory
scalability	<i>harder</i>	<i>easier</i>
programmability	<i>easier</i>	<i>harder</i>

- ▶ shared memory permits parallelizing serial program gradually, focusing on worst bottlenecks first
- ▶ distributed memory requires partitioning and distributing both data and work across processors, which usually rules out incremental parallelization

Categories of parallel architectures

- ▶ vector or array processor
- ▶ SMP: symmetric multiprocessor
- ▶ MPP: massively parallel multiprocessor
- ▶ DSM: distributed shared memory
- ▶ clusters
- ▶ hybrids
 - ▶ SMP or MPP with vector processors
 - ▶ networked clusters of SMPs
 - ▶ SMP+GPGPU

Generic parallel architecture



- ▶ a trivial but powerful observation: access to memory implies access to information
- ▶ hence, it becomes a determining factor for both hardware and algorithm design

Memory hierarchy


- ▶ high performance architectures have a multi-tier memory hierarchy
 - ▶ registers
 - ▶ on-chip caches, usually referred to as level 1
 - ▶ off-chip caches (level 2)
 - ▶ random access memory
 - ▶ remote memory (off processor)
 - ▶ virtual memory, known as paging memory, that usually involves secondary storage
 - ▶ secondary storage (disks)
 - ▶ tertiary storage (tapes)
- ▶ these have latencies and bandwidths that vary by orders of magnitude
- ▶ *cache misses* are the most frequently cited reason why real codes only achieve a small fraction of the benchmarked performance of a CPU
 - ▶ really small: 10% of peak is considered a success!

Parallel programming paradigms

- ▶ *functional languages*: specify what to compute, not how
- ▶ *parallelizing compilers*: automatic or semi-automatic detection of parallelism in serial code; mostly loop-unrolling, often with the help of special mark up (pragmas)
- ▶ *object oriented*: parallelism encapsulated within distributed objects
- ▶ *data parallel*: simultaneous operations on memory; mostly arrays
- ▶ *shared memory*: multiple threads executing a pool of tasks using common memory
- ▶ *remote memory access*: one sided put/get communication between processes
- ▶ *message passing*: two sided, coördinated send/receive communication between processes

Designing parallel algorithms

- ▶ *identification*: identify the part of the problem that can be parallelized
- ▶ *partition*: decompose the parallelizable part into fine-grained tasks
- ▶ *communication*: determine the necessary communication patterns among tasks
- ▶ *coarsening*: combine into coarser tasks and adjust the communication patterns
- ▶ *task mapping*: assign tasks to processors



Paradigms for parallel algorithms

- ▶ *embarrassingly parallel*: mostly independent tasks
- ▶ *functional decomposition*: based on computational task (activity)
- ▶ *data parallel*: aka *loop-level* parallelism: array operations
- ▶ *domain decomposition*: based on the distribution of data
- ▶ *divide-and-conquer*: tree-like partitioning
- ▶ *pipelining*: multiple overlapping stages

Communication issues

- ▶ latency and bandwidth
- ▶ routing and switching – not much of an issue any more
- ▶ contention, flow control and aggregate bandwidth
- ▶ collective communication
 - ▶ one to many: broadcast, scatter
 - ▶ many to one: gather, reduction, scan
 - ▶ all to all
 - ▶ synchronization barrier
- ▶ assigning work to processors
 - ▶ partitioning
 - ▶ granularity
 - ▶ mapping
 - ▶ scheduling
 - ▶ load balancing

Factors determining performance

- ▶ *concurrency*: maximize the work that can be done in parallel
- ▶ *load balance*: make sure the load is (and stays) divided evenly
- ▶ *parallel overhead*: work not present in the equivalent serial computation
 - ▶ process startup and shutdown costs
 - ▶ communication
 - ▶ synchronization
 - ▶ redundancy
 - ▶ speculative work

Computational models

- ▶ abstractions for architecture, algorithm analysis, and performance modeling
 - ▶ PRAM: Parallel random access machine
 - ▶ LogP: latency, overhead, gap, processors
 - ▶ BSP: bulk synchronous parallel
 - ▶ CSP: communicating sequential processes
 - ▶ ... and many others
- ▶ all occasionally useful tools for reasoning about implementation strategies for real programs, since they steer you away from common but perhaps non-obvious mistakes
- ▶ unfortunately, none is a substitute for understanding the characteristics of your target platform