ACM/CS 114
Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Winter 2012

Hello world

```
1  #include <pthread.h>
2  #include <stdio.h>
3  #define THREADS 10
4
5  void* hello(void* threadID) {
6     long id = (long) threadID;
7     printf("hello from %02ld/%0d\n", id, THREADS);
8     pthread_exit(NULL);
9     return NULL;
10 }
11
12 int main(int argc, char* argv[]) {
13    long id;
14    int status;
15    pthread_t threads[THREADS];
16
17    for (id=0; id<THREADS; id++) {
18       printf("creating thread %02ld\n", id);
19       status = pthread_create(&threads[id], NULL, hello, (void*) id);
20       if (status) {
21          printf("error %d in pthread_create\n", status);
22       }
23    }
24    /* there is a problem here... */
25    pthread_exit(NULL);
26    return 0;
27 }
```

Joining and detaching

- in the example in Slide 2, the main thread exits without knowing whether any of the threads it spawned have finished
  - saying "hello" is asynchronous
  - but gathering the results of parallel calculations normally isn't
- *thread synchronization* can be achieved using pthread_join
  - the pthread_create caller saves the thread id
  - the thread is scheduled, executes, and calls pthread_exit
  - any other thread can wait for this thread to finish by calling pthread_join with the saved thread id and also retrieve the termination status
- for this to work, a thread must be *joinable*
  - controlled by the thread creation attributes
  - for portability, you should always mark your joinable threads explicitly
- a thread that will never be joined may be *detached*
  - by setting the corresponding attribute during thread creation
  - or, by calling pthread_detach at any point
  - detaching a thread saves some system resources

Creating mutexes

- a *mutex* is a locking mechanism that helps guarantee exclusive access to a section of code, most often to control access to shared variables
- mutexes are created using

```
1 int pthread_mutex_init(
2   pthread_mutex_t* mutex, const pthread_mutexattr_t* attr);
```

  - they start out unlocked
  - the `attr` enables more advanced (but perhaps non-portable) use
- mutexes are destroyed using

```
1 int pthread_mutex_destroy(pthread_mutex_t* mutex);
```

 - destroy mutexes you are no longer using to prevent resource leakage

Locking and unlocking mutexes

- ▶ threads manipulate mutexes through

```
1 int pthread_mutex_lock(pthread_mutex_t* mutex);
2 int pthread_mutex_trylock(pthread_mutex_t* mutex);
3 int pthread_mutex_unlock(pthread_mutex_t* mutex);
```

- ▶ pthread_mutex_lock attempts to gain exclusive access
  - ▶ if the mutex is unlocked, it locks it and returns
  - ▶ otherwise, it blocks until the mutex is unlocked; when the mutex is unlocked, it locks it and returns
- ▶ pthread_mutex_unlock attempts to release a mutex
  - ▶ if it was previously locked by this thread, the mutex is unlocked
  - ▶ if it was not previously locked, the call returns with an error code
  - ▶ if it was locked, but not by the calling thread, the call returns an error code
- ▶ pthread_mutex_trylock attempts to lock the mutex
  - ▶ if it is unlocked, the call locks it and returns
  - ▶ if it is locked, the call returns immediately with a *busy* error code
- ▶ locking and unlocking mutexes is explicitly orchestrated by the programmer
- ▶ when multiple threads are blocked waiting for a mutex, there is no way to predict which one will succeed when the mutex becomes available

Reduction on a shared memory machine

Example reduction using threads

```
1   #include <pthread.h>
2   #include <stdio.h>
3   #define THREADS 10
4   /* global variables -- yuck! */
5   long sum = 0;
6   pthread_mutex_t mutex;
7   /* worker */
8   void* squares(void* threadID) {
9       long id = (long) threadID;
10      pthread_mutex_lock(&mutex);
11      sum += id*id;
12      pthread_mutex_unlock(&mutex);
13      pthread_exit(NULL);
14      return NULL;
15  }
16  /* main program */
17  int main(int argc, char* argv[]) {
18      long id;
19      pthread_t threads[THREADS];
20      pthread_mutex_init(&mutex, NULL);
21      /* create some threads */
22      for (id=0; id<THREADS; id++) {
23          pthread_create(&threads[id], NULL, squares, (void*) id);
24      }
25      /* wait for them to finish */
26      for (id=0; id<THREADS; id++) {
27          pthread_join(threads[id], NULL);
28      }
29      /* print the result */
30      printf("sum = %ld\n", sum);
31      /* exit */
32      pthread_exit(NULL);
33      return 0;
34  }
```

Condition variables

- condition variables build upon mutexes to enable threads to signal each other when some condition is met

- they are created using

```
int pthread_cond_init(
  pthread_cond_t* condition, const pthread_condattr_t* attr);
```

- and destroyed using

```
int pthread_cond_destroy(pthread_cond_t* condition);
```

Using condition variables

- the following three routines implement the condition variable semantics

```
1 int pthread_cond_wait(pthread_cond_t* cond, pthread_mutex_t* mutex);
2 int pthread_cond_signal(pthread_cond_t* cond);
3 int pthread_cond_broadcast(pthread_cond_t* cond);
```

- `pthread_cond_wait` blocks the calling thread until the specified condition is *signaled*
  - it must be called with the mutex locked by the calling thread
  - `pthread_cond_wait` releases the lock while the thread is blocked
  - after the matching signal is received, the thread is awakened and the mutex locked
  - the thread is responsible for releasing the mutex when it is done
- `pthread_cond_signal` wakes up a thread that is waiting for the given condition variable
  - mutex must be locked before calling it
  - mutex must be unlocked after signaling, so blocking threads can be awakened
- `pthread_cond_broadcast` can be used instead if multiple threads are waiting for a signal

Condition variable caveats

- be careful with condition variables; make sure that
    - a thread has called `pthread_cond_wait` before any thread calls `pthread_cond_signal`
    - the mutex associated with the condition is locked before calling `pthread_cond_wait`, otherwise it might *not block*
    - the thread that calls `pthread_cond_signal` unlocks the associated mutex, otherwise the threads waiting for the signal will continue to block

Attributes of threads, mutexes and condition variables

- ▶ threads, mutexes and condition variables have associated attribute structures that can be used to tune the default creation parameters
- ▶ they are created and destroyed using

```
1 int pthread_attr_init(pthread_attr_t* attr);
2 int pthread_attr_destroy(pthread_attr_t* attr);
3
4 int pthread_mutexattr_init(pthread_mutexattr_t* attr);
5 int pthread_mutexattr_destroy(pthread_mutexattr_t* attr);
6
7 int pthread_condattr_init(pthread_condattr_t* attr);
8 int pthread_condattr_destroy(pthread_condattr_t* attr);
```

- ▶ typically, the defaults are adequate and tuned to the details of the operating system
- ▶ if you make excessive use of the stack, e.g. large arrays as local variable or deep recursion, you might want to know about

```
1 int pthread_attr_getstacksize(pthread_attr_t* attr, size_t* size);
2 int pthread_attr_setstacksize(pthread_attr_t* attr, size_t size);
```

Other useful routines

- a thread can access its unique id assigned by the system by calling

```
pthread_t pthread_self(void);
```

- since system thread ids are opaque types, you cannot use == to compare them; instead, use

```
int pthread_equal(pthread_t id1, pthread_t id2);
```

- you can place all thread initialization code in a startup routine and call

```
int pthread_once(pthread_once_t* control_structure, void (*startup_routine)
```

Advanced topics

- there is quite a bit more in the standard
- *keys*: creating and accessing per-thread data
    - as the code get more complicated, it becomes increasingly difficult to pass complete thread-specific information from function to function
    - possible solutions:
        - the FORTRAN syndrome, where subroutines end up having dozens of arguments
        - global variables
        - an associative container that allows each thread to store and retrieve arbitrary data
- finer control over thread scheduling
    - scheduling algorithms and priorities are implementation dependent
    - there are routines in the standard that enable explicit tuning
    - the standard guarantees that the routines will be *available*, but they don't have to be *implemented*
- condition variable sharing across processes
- explicitly canceling threads
- the somewhat complicated interactions between threads and signals
- other synchronization constructs: barriers and read/write locks

Summary

- well-designed threaded programs must follow the same strategy as any other concurrent program

  - identify the work that can be done concurrently
  - partition it in terms of work units, the fine grain tasks
  - analyze the communication patterns among work units with an eye for critical sections and protecting shared data structures
  - coarsen into threads, define the mutex categories and synchronization points
  - let the OS schedule the threads onto physical processors
- debugging threaded programs is very difficult
  - preventing bugs through careful design is critical
  - so is instrumenting the program to gain confidence in its execution