

ACM/CS 114

Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Winter 2012

Messages

- ▶ in general, data exchanges through MPI calls involve
 - ▶ a communicator
 - ▶ specifies which processes participate in the exchange
 - ▶ resolves process ranks into processes
 - ▶ *collective* operations involve the entire communicator
 - ▶ *point-to-point* operations require the rank of the message source or destination
 - ▶ the details of the message payload
 - ▶ the address of the source buffer
 - ▶ the data type of the buffer contents
 - ▶ the number of items in the buffer
- ▶ MPI provides some data abstractions to
 - ▶ hide machine dependencies in the data representations to enhance portability and support heterogeneous clusters
 - ▶ support user defined data types
 - ▶ support non-contiguous data layouts

Collective operations: global reductions

- ▶ *collective* operations involve all processes in a given communicator
- ▶ the MPI version of our global reduction example uses

```
1 int MPI_Allreduce(  
2     void* send_buffer, void* recv_buffer,  
3     int count, MPI_Datatype datatype, MPI_Op operation,  
4     MPI_Comm communicator  
5 );
```

- ▶ example legal values for MPI_Datatype
 - ▶ C: MPI_INT, MPI_LONG, MPI_DOUBLE
 - ▶ FORTRAN: MPI_INTEGER, MPI_DOUBLE_PRECISION, MPI_COMPLEX
- ▶ legal values for MPI_Op
 - ▶ MPI_MAX, MPI_MIN, MPI_MAXLOC, MPI_MINLOC
 - ▶ MPI_SUM, MPI_PROD
 - ▶ MPI_LAND, MPI_LOR, MPI_LXOR
 - ▶ MPI_BAND, MPI_BOR, MPI_BXOR
 - ▶ MPI_REPLACE

Example reduction using MPI

```
1 #include <mpi.h>
2 #include <stdio.h>
3
4 int main(int argc, char* argv[]) {
5     int status;
6     int rank;
7     int square, sum;
8
9     /* initialize MPI */
10    status = MPI_Init(&argc, &argv);
11    if (status != MPI_SUCCESS) {
12        printf("error in MPI_Init; aborting...\n");
13        return status;
14    }
15
16    /* get the process rank */
17    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
18    /* form the square */
19    square = rank*rank;
20    /* each process contributes the square of its rank */
21    MPI_Allreduce(&square, &sum, 1, MPI_INT, MPI_SUM, MPI_COMM_WORLD);
22    /* print out the result */
23    printf("%03d: sum = %d\n", rank, sum);
24
25    /* shut down MPI */
26    MPI_Finalize();
27
28    return 0;
29 }
```