

ACM/CS 114

Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Winter 2012

- ▶ the function

```
1 double MPI_Wtime();
```

returns the time in seconds from some arbitrary time in the past

- ▶ guaranteed not to change only for the duration of the process
- ▶ you can compute the elapsed time for any program segment by making calls at the beginning and the end and computing the difference
- ▶ no guarantees about synchronized clocks among different processes
- ▶ you can compute the clock resolution by using

```
1 double MPI_Wtick();
```

Other collective operations

- ▶ `MPI_Scan` computes partial reductions: the p^{th} process receives the result from processes 0 through $p - 1$

```
1 int MPI_Scan(  
2     void* send_buffer, void* recv_buffer,  
3     int count, MPI_Datatype datatype, MPI_Op operation,  
4     MPI_Comm communicator  
5 );
```

- ▶ `MPI_Reduce` collects the result at only the given process root

```
1 int MPI_Reduce(  
2     void* send_buffer, void* recv_buffer,  
3     int count, MPI_Datatype datatype, MPI_Op operation,  
4     int root, MPI_Comm communicator  
5 );
```

- ▶ synchronization is also a global operation:


```
1 int MPI_Barrier(MPI_Comm communicator);
```

participating processes block at a barrier until they have all reached it

Scatter

- ▶ `MPI_Scatter` sends data from `root` to all processes

```
1 int MPI_Scatter(  
2     void* send_buffer, int send_count, MPI_Datatype send_datatype,  
3     void* recv_buffer, int recv_count, MPI_Datatype recv_datatype,  
4     int root, MPI_Comm communicator  
5 );
```




- ▶ it is as if the data in `send_buffer` were split in p segments, and the i^{th} process receives the i^{th} segment
- ▶ the `send_xxx` arguments are only meaningful for `root`; they are ignored for other processes
- ▶ the arguments `root` and `communicator` must be passed identical values by all processes

Gather

- ▶ the converse is `MPI_Gather` with `root` receiving data from all processes

```
1 int MPI_Gather(  
2     void* send_buffer, int send_count, MPI_Datatype send_datatype,  
3     void* recv_buffer, int recv_count, MPI_Datatype recv_datatype,  
4     int root, MPI_Comm communicator  
5 );
```



- ▶ it is as if p messages, one from each processes, were concatenated in rank order and placed at `recv_buffer`
- ▶ the `recv_xxx` arguments are only meaningful for `root`; they are ignored for other processes
- ▶ the arguments `root` and `communicator` must be passed identical values by all processes

Broadcasting operations

- ▶ `MPI_Alltoall` sends data from all processes to all processes in a global scatter/gather

```
1 int MPI_Alltoall(  
2     void* send_buffer, int send_count, MPI_Datatype send_datatype,  
3     void* recv_buffer, int recv_count, MPI_Datatype recv_datatype,  
4     MPI_Comm communicator  
5 );
```

- ▶ use `MPI_Bcast` to send the contents of a buffer from `root` to all processes in a communicator

```
1 int MPI_Bcast(  
2     void* buffer, int count, MPI_Datatype datatype,  
3     int root, MPI_Comm communicator  
4 );
```

Data movement patterns for the collective operations

