

ACM/CS 114

Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Winter 2012

Approximating Li_2 using a numerical quadrature

- ▶ the second homework assignment involved $\text{Li}_2(z)$, defined by


$$\text{Li}_2(z) := - \int_0^z dz' \frac{\log(1 - z')}{z'}$$

- ▶ the assignment asked for approximating this integral using a simple quadrature based on the mid-point rule

$$\text{Li}_2(z) \approx \text{Li}_2(z, N) := - \frac{z}{N} \sum_{n=0}^{N-1} \frac{\log(1 - z')}{z'} \Big|_{z' = (n + \frac{1}{2}) \frac{z}{N}}$$

Quadrature rule

$$\text{Li}_2(z) \approx -dz \sum_{n=0}^{N-1} \frac{\log(1 - z')}{z'} \bigg|_{z'=(n+\frac{1}{2})dz}$$



Implementations

- ▶ three implementations
 - ▶ sequential: to get a feeling for how to convert the algorithm into a functioning program
 - ▶ parallel using threads: to walk through the parallelization steps and use `pthread`s to get better performance
 - ▶ parallel using MPI: to get a feel for how MPI-based programs solve the task partitioning problem
- ▶ let's walk through composing, building and running
 - ▶ on my desktop, and on `shc.cacr.caltech.edu`

Sequential implementation - part 1

► the preamble

```
1 #include <getopt.h> // for getopt and friends
2 #include <cstdlib> // for atof
3 #include <cmath> // for the correct abs, log
4
5 #include <map>
6 #include <iostream>
7 #include <iomanip>
```

► quadrature using the midpoint rule to avoid the singularities

```
8 // dilog
9 double dilog(double z, long N) {
10     // initialize
11     double dx = z/N;
12     double x = dx/2;
13     double sum = 0;
14     // loop
15     for (long i=0; i < N; i++) {
16         sum += std::log(1-x)/x;
17         x += dx;
18     }
19     // return; don't forget the sign
20     return -dx * sum;
21 }
```

Sequential implementation - part 2

- using the command line to set z and the number of subdivisions N

```
23 // main program
24 int main(int argc, char* argv[]) {
25     // default values for the command line options
26     long N = 1000;
27     double z = 1.0;
28
29     // read the command line
30     int command;
31     while ((command = getopt(argc, argv, "z:N:")) != -1) {
32         switch (command) {
33             // get the argument of the dilogarithm
34             case 'z':
35                 z = atof(optarg);
36                 break;
37             // get the number of subdivisions
38             case 'N':
39                 N = (long) atof(optarg);
40                 break;
41         }
42     }
```

Sequential implementation - part 3

► error checking and computation of the numerical integral

```
43 // error checking
44 // abort if N < 1
45 if (N < 1) {
46     std::cout
47         << "the number of subdivisions must be positive"
48         << std::endl;
49     return 0;
50 }
51
52 // abort for z > 1 to avoid dealing with the imaginary part
53 if (z > 1.0) {
54     std::cout << "math domain error: z > 1" << std::endl;
55     return 0;
56 }
57
58 // compute
59 double value = dilog(z, N);
```

Sequential implementation - part 4

- computing the error and printing out the results

```
60 // build a naive database of the known dilogarithm values
61 const double pi = M_PI;
62 std::map<double, double> answers;
63 answers[1.0] = pi*pi/6;
64 answers[-1.0] = -pi*pi/12;
65
66 // print out the value
67 std::cout << "Li2(" << z << ")="
68     << std::setprecision(17) << std::endl
69     << " computed: " << value << std::endl;
70 // check whether we know the right answer
71 std::map<double,double>::const_iterator lookup = answers.find(z);
72 if (lookup != answers.end()) {
73     // and if we do, print it out
74     double exact = lookup->second;
75     std::cout << " exact: " << exact << std::endl;
76     // compute the approximation error and print it out
77     double error = std::abs(exact-value)/exact;
78     std::cout
79         << std::setiosflags(std::ios_base::scientific)
80         << " error: " << error << std::endl;
81 }
82
83 return 0;
84 }
```


Building and running the sequential driver

```
1 #> g++ dilog-quadrature_sequential
2 #> dilog-sequential -N 1e9 -z 1.0
3 Li2(1)=
4   computed: 1.6449339414016682
5   exact: 1.6449340668482264
6   error: 7.62623625958871898e-08
7
8 real    0m19.885s
9 user    0m19.877s
10 sys     0m0.003s
11 #>
```

Threaded implementation - part 1

► the preamble

```
1 #include <getopt.h> // for getopt and friends
2 #include <pthread.h>
3
4 #include <stdio>
5 #include <stdlib> // for atof
6 #include <cmath>
7
8 #include <map>
9 #include <iostream>
10 #include <iomanip>
```

Threaded implementation - part 2

► private and shared data structures

```
12 // shared information
13 struct problem {
14     int workers;      // total number of threads
15     double dz;        // the width of each subdivision
16     double sum;       // storage for the partial computations
17
18     pthread_mutex_t lock; // mutex to control access to the sum
19 };
20
21 // thread specific information
22 struct context {
23     // thread info
24     int id;
25     pthread_t descriptor;
26     // the workload for this thread
27     long subdivisions; // number of subdivisions
28     double z_low;      // the lower limit of integration
29     double partial;    // record the partial sum computed by this thread
30     // the shared problem information
31     problem* info;
32 };
```

Threaded implementation - part 3

► the coarse grain task

```
33 // worker
34 void* worker(void* arg) {
35     context* ctxt = (context *) arg;
36     // pull the problem information from the thread context
37     double dz = ctxt->info->dz;
38     double z = ctxt->z_low + dz/2;
39     // loop over the subdivisions assigned to this thread
40     double sum = 0.0;
41     for (long i=0; i < ctxt->subdivisions; i++) {
42         sum += std::log(1-z)/z;
43         z += dz;
44     }
45     // multiply by the width of each subdivision and adjust the sign
46     sum *= -dz;
47
48     // grab the lock
49     pthread_mutex_lock(&(ctxt->info->lock));
50     // store the result
51     ctxt->info->sum += sum;
52     // and release the lock
53     pthread_mutex_unlock(&(ctxt->info->lock));
54
55     // all done
56     return 0;
57 }
```