

# ACM/CS 114

## Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Winter 2012



# An example

- specifically,
  - let  $\Omega$  be the unit box in two dimensions
  - and let  $\phi$  satisfy the following boundary conditions

$$\begin{aligned}\phi(x, 0) &= \sin(\pi x) & 0 \leq x \leq 1 \\ \phi(x, 1) &= e^{-\pi} \sin(\pi x) & 0 \leq x \leq 1 \\ \phi(0, y) &= \phi(1, y) = 0 & 0 \leq y \leq 1\end{aligned}\tag{5}$$

- the exact solution is given by


$$\phi(x, y) = e^{-\pi y} \sin(\pi x)\tag{6}$$

- we will solve this equation using the Jacobi iterative scheme:
  - make an initial guess for  $\phi$  over a discretization of  $\Omega$
  - apply the boundary conditions
  - interpret Eq. 4 as an update step to compute the next iteration

$$\begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array}_t = \frac{1}{4} \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array}_{t-1}\tag{7}$$

- stop when a convergence criterion is met

# The solution



# Implementation strategy

- ▶ grid resolution:
  - ▶ ideally determined by analyzing the boundary conditions, since discrete sampling may wash out sharp features
  - ▶ for our simple example, this can be done as part of the solver initialization
  - ▶ we will use an  $N \times N$  grid and let  $N$  be user specified so we can control the problem size

$$\delta_x = \delta_y = \frac{1}{N-1} \quad (8)$$


- ▶ data layout
  - ▶ investigate the effect of data locality by trying out various layouts
- ▶ setting up the update
  - ▶ we only need to keep track of two iterants
  - ▶ can be done in place; do you see how?
- ▶ convergence criterion
  - ▶ we will stop iterating when

$$\max_{\Omega} (\phi_t - \phi_{t-1}) < \epsilon \quad (9)$$

and let the user specify  $\epsilon$

# Parallelization

- ▶ the finest grain of work is clearly the cell update based on the value of its four nearest neighbors
- ▶ the shared memory implementation requires
  - ▶ a scheme so that threads can update cells without the need for locks
  - ▶ while maximizing locality of data access
  - ▶ even the computation of the convergence criterion can be parallelized
- ▶ with MPI
  - ▶ must partition the mesh among processes
  - ▶ each process work on its own subgrid
  - ▶ communication is required every iteration
  - ▶ parallel convergence testing involves a collective operation



# Sequential implementation - user interface

```
77 // main program
78 int main(int argc, char* argv[]) {
79     // default values for our user configurable settings
80     size_t N = 10;
81     double tolerance = 1.0e-6;
82     const char* filename = "laplace.csv";
83
84     // read the command line
85     int command;
86     while ((command = getopt(argc, argv, "N:e:o:")) != -1) {
87         switch (command) {
88             // get the convergence tolerance
89             case 'e':
90                 tolerance = atof(optarg);
91                 break;
92             // get the grid size
93             case 'N':
94                 N = (size_t) atof(optarg);
95                 break;
96             // get the name of the output file
97             case 'o':
98                 filename = optarg;
99         }
100     }
```

# Sequential implementation - driving the solver

```
102 // allocate space for the solution
103 Grid potential(N);
104
105 // initialize and apply our boundary conditions
106 initialize(potential);
107
108 // call the solver
109 laplace(potential, tolerance);
110
111 // open a stream to hold the answer
112 std::fstream output(filename, std::ios_base::out);
113
114 // build a visualizer and render the solution in our chosen format
115 Visualizer visualizer;
116 visualizer.csv(potential, output);
117
118 // all done
119 return 0;
120 }
```



# Sequential implementation - the preamble

- back up to the beginning of the file

```
1 #include <getopt.h>
2 #include <cmath>
3 #include <cstdlib>
4 #include <fstream>
5 #include <iostream>
6
7 // forward declarations
8 class Grid;
9 class Visualizer;
10
11 // the solver; does nothing for the time being
12 void initialize(Grid & grid) {};
13 void laplace(Grid & grid, double tolerance) {};
```

- we have separated out *visualization* in a different object to support different formats without disturbing the data representation
- `initialize` and `laplace` have trivial implementations for now
  - enables testing the scaffolding without worrying about the solver implementation just yet

# Sequential implementation - the grid object stub

```
15 // the solution representation
16 class Grid {
17     // interface: TBD
18 public:
19
20     // meta methods
21 public:
22     Grid(size_t size);
23     ~Grid();
24
25     // private data members: TBD
26 private:
27
28     // disabled interface
29     // grid will own dynamic memory, so don't let the compiler screw up
30 private:
31     Grid(const Grid &);
32     const Grid & operator= (const Grid &);
33 };
34
35 // the grid implementation
36 Grid::Grid(size_t size) {
37 }
38
39 Grid::~~Grid() {
40 }
```

# Sequential implementation - the visualizer stub

```
97 // the visualizer class
98 class Visualizer {
99     // local type aliases
100 public:
101     typedef std::ostream stream_t;
102
103     // interface
104 public:
105     void csv(const Grid & grid, stream_t & stream);
106
107     // meta methods
108 public:
109     inline Visualizer() {}
110 };
111
112 // the Visualizer class implementation
113 void Visualizer::csv(const Grid & grid, Visualizer::stream_t & stream) {
114     return;
115 }
```

- ▶ the code now compiles and links
  - ▶ consistency check that the object collaborations are ok, for now
  - ▶ can be tested for command line option parsing