

ACM/CS 114

Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Winter 2012

Shortcomings

- ▶ numerics:
 - ▶ it converges very slowly; other update *schemes* improve on this
 - ▶ our approximation is very low order, so it takes very large grids to produce a few digits of accuracy
 - ▶ the convergence criterion has some unwanted properties; it triggers
 - ▶ prematurely: large swaths of constant values may never get updated
 - ▶ it would trigger even if we were updating the wrong grid!
- ▶ design:
 - ▶ separate the problem specification from its solution
 - ▶ there are other objects lurking, waiting to be uncovered
 - ▶ someone should make the graphic visualizer
 - ▶ restarts anybody?
 - ▶ how would you try out different convergence criteria? update schemes? memory layouts?
- ▶ usability:
 - ▶ supporting interchangeable parts requires damage to the top level driver
 - ▶ to enable the user to make the selection
 - ▶ to expose new command line arguments that configure the new parts

Assessing our fundamentals

- ▶ `Grid` is a good starting point for abstracting structured grids
 - ▶ assumes ownership of the memory associated with a structured grid
 - ▶ encapsulates the indexing function
 - ▶ extend it to
 - ▶ support different memory layout strategies
 - ▶ support non-square grids (?)
 - ▶ support non-uniform grids (?)
 - ▶ higher dimensions
 - ▶ if you need any of these, consider using one of the many excellent class libraries written by experts
- ▶ `Visualizer`, under another name, can form the basis for a more general persistence library
 - ▶ to support HDF5, NetCDF, bitmaps, voxels, etc.

The Problem class: the interface

```
1 // the solution representation
2 class acml14::laplace::Problem {
3     //typedefs
4 public:
5     typedef std::string string_t;
6     // interface
7 public:
8     inline string_t name() const;
9     inline const Grid & exact() const;
10    inline const Grid & deviation() const;
11    inline Grid & solution();
12    inline const Grid & solution() const;
13    inline Grid & error();
14    inline const Grid & error() const;
15    // abstract
16    virtual void initialize() = 0;
17    virtual void initialize(Grid &) const = 0;
18    // meta methods
19 public:
20    inline Problem(string_t name, double width, size_t points);
21    virtual ~Problem();
22    // data members
```

The Problem class: the data

```
23 protected:
24     string_t _name;
25     double _delta;
26     Grid _solution;
27     Grid _exact;
28     Grid _error;
29     Grid _deviation;
30     // disable these
31 private:
32     Problem(const Problem &);
33     const Problem & operator= (const Problem &);
34 };
```

The Example class

```
1 class acm114::laplace::Example : public acm114::laplace::Problem {
2     // interface
3 public:
4     virtual void initialize();
5     virtual void initialize(Grid &) const;
6
7     // meta methods
8 public:
9     inline Example(string_t name, double width, size_t points);
10    virtual ~Example();
11
12    // disable these
13 private:
14     Example(const Example &);
15     const Example & operator= (const Example &);
16 };
```

The Solver class

```
1 class acm114::laplace::Solver {
2     // interface
3 public:
4     virtual void solve(Problem &) = 0;
5
6     // meta methods
7 public:
8     inline Solver();
9     virtual ~Solver();
10
11     // data members
12 private:
13
14     // disable these
15 private:
16     Solver(const Solver &);
17     const Solver & operator= (const Solver &);
18 };
```

The Jacobi class

```
1 class acml14::laplace::Jacobi : public acml14::laplace::Solver {
2     // interface
3 public:
4     virtual void solve(Problem &);
5
6     // meta methods
7 public:
8     inline Jacobi(double tolerance, size_t workers);
9     virtual ~Jacobi();
10
11     // implementation details
12 protected:
13     virtual void _solve(Problem &);
14     static void * _update(void *);
15
16     // data members
17 private:
18     double _tolerance;
19     size_t _workers;
20
21     // disable these
22 private:
23     Jacobi(const Jacobi &);
24     const Jacobi & operator= (const Jacobi &);
25 };
```


Parallelization using threads

- ▶ the shared memory implementation requires
 - ▶ a scheme so that threads can update cells without the need for locks
 - ▶ while maximizing locality of data access
 - ▶ even the computation of the convergence criterion can be parallelized
- ▶ parallelization strategy
 - ▶ we will focus on parallelizing the iterative grid update
 - ▶ grid initialization, visualization, computing the exact answer and the error field do not depend on the *number of iterations*
 - ▶ the finest grain of work is clearly an individual cell update based on the value of its four nearest neighbors
 - ▶ for this two dimensional example, we can build coarser grain tasks using
 - ▶ horizontal or vertical strips
 - ▶ non-overlapping blocks
 - ▶ the strategy gets more complicated if you want to perform the update in place
 - ▶ the communication patterns are trivial for the double buffering layout; only the final update of the convergence criterion requires any locking
 - ▶ each coarse grain task can be assigned to a thread

Required changes to the sequential solution

- ▶ what is needed
 - ▶ an object to hold the problem information shared among the threads
 - ▶ the per-thread administrative data structure that holds the thread id and the pointer to the shared information
 - ▶ this is the argument to `pthread_create`
 - ▶ a mutex to protect the update of the global convergence criterion
 - ▶ a `pthread_create` compatible worker routine
 - ▶ a change at the top-level driver to enable the user to choose the number of threads
- ▶ and a strategy for managing the thread life cycle
 - ▶ synchronization is trivial if
 - ▶ we spawn our threads to perform the updates of a single iteration
 - ▶ harvest them
 - ▶ check the convergence criterion
 - ▶ stop, or respawn them if another iteration is necessary
 - ▶ can the convergence test be done in parallel?
 - ▶ so we don't have to pay the create/harvest overhead?
 - ▶ if so, how do we guarantee correctness and consistency?

Threaded Jacobi: thread data

```
1 struct Task {
2     // shared information
3     size_t workers;
4     Grid & current;
5     Grid & next;
6     double maxDeviation;
7     // mutex to control access to the convergence criterion
8     pthread_mutex_t lock;
9
10    // constructor
11    Task(size_t workers, Grid & current, Grid & next) :
12        workers(workers), current(current), next(next), maxDeviation(0.0) {
13        pthread_mutex_init(&lock, 0);
14    }
15    // destructor
16    ~Task() {
17        pthread_mutex_destroy(&lock);
18    }
19 };
20
21 struct Context {
22     // thread info
23     size_t id;
24     pthread_t descriptor;
25     Task * task;
26 };
```

Threaded Jacobi: driving the update

```
28 void Jacobi::solve(Problem & problem) {
29     // initialize the problem
30     problem.initialize();
31     // do the actual solve
32     _solve(problem);
33     // compute and store the error
34     std::cout << " computing absolute error" << std::endl;
35     // compute the relative error
36     Grid & error = problem.error();
37     const Grid & exact = problem.exact();
38     const Grid & solution = problem.solution();
39
40     for (size_t j=0; j < exact.size(); j++) {
41         for (size_t i=0; i < exact.size(); i++) {
42             if (exact(i,j) == 0.0) {
43                 error(i,j) = std::abs(solution(i,j));
44             } else {
45                 error(i,j) = std::abs(solution(i,j) - exact(i,j))/exact(i,j);
46             }
47         }
48     }
49     std::cout << " --- done." << std::endl;
50     return;
51 }
```