

# ACM/CS 114

## Parallel algorithms for scientific applications

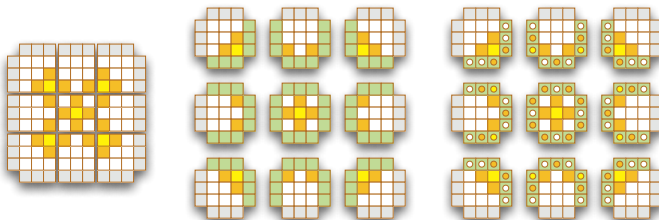
Michael A. G. Aïvázis

California Institute of Technology

Winter 2012

# Parallelization with MPI

- ▶ the MPI implementation will require careful data management
  - ▶ we must partition the mesh among processes
  - ▶ each process work on its own subgrid
    - ▶ it will allocate its own memory, for both actual data and the guard zones
    - ▶ it must locate its patch in physical space
  - ▶ communication is required every iteration
    - ▶ so that neighbors can synchronize their boundaries
    - ▶ think of the synchronization as a kind of boundary condition!
  - ▶ parallel convergence testing involves a collective operation



# A little bit of help

- ▶ MPI supports this common use case through a Cartesian *virtual topology*
  - ▶ a special communicator with a map from a  $d$ -dimensional virtual process grid to the normal linear process ranks
  - ▶ and local operations that enable you to discover the ranks of your virtual neighbors
  - ▶ there is even a special form of send/receive so that you don't have to worry about contention and race conditions during the boundary synchronization
- ▶ to create a Cartesian communicator

```
1 int MPI_Cart_create(MPI_Comm oldcomm,  
2     int ndims, int* layout, int* periods, int reorder, MPI_Comm* newcomm)
```

- ▶ to find out the coordinates of a process in the virtual grid given its rank

```
1 int MPI_Cart_coords(MPI_Comm cartesian,  
2     int rank, int ndims, int* coords);
```

- ▶ you can also find out the ranks of your neighbors

```
1 int MPI_Cart_shift(MPI_Comm cartesian,  
2     int dimension, int shift, int* origin, int* neighbor);
```

# The MPI driver, part 1

```
26 int main(int argc, char* argv[]) {
27     int status;
28     // initialize mpi
29     status = MPI_Init(&argc, &argv);
30     if (status) {
31         throw("error in MPI_Init");
32     }
33     // get my rank in the world communicator
34     int worldRank, worldSize;
35     MPI_Comm_rank(MPI_COMM_WORLD, &worldRank);
36     MPI_Comm_size(MPI_COMM_WORLD, &worldSize);
37     size_t processors = static_cast<size_t>(std::sqrt(worldSize));
38
39     // default values for our user configurable settings
40     size_t n = 9; // points per processor
41     size_t threads = 1;
42     double tolerance = 1.0e-3;
43
44     // read the command line
45     int command;
46     while ((command = getopt(argc, argv, "n:e:t:")) != -1) {
47         switch (command) {
48             // get the convergence tolerance
49             case 'e':
50                 tolerance = atof(optarg);
51                 break;
52             // get the grid size
53             case 'n':
54                 n = (size_t) atof(optarg);
55                 break;
56             // get the number of threads
57             case 't':
58                 threads = (size_t) atoi(optarg);
59                 break;
60         }
61     }
```

# The MPI driver, part 2

```
62 // print out the chosen options
63 if (worldRank == 0) {
64     for (int arg = 0; arg < argc; ++arg) {
65         std::cout << argv[arg] << " ";
66     }
67     std::cout
68         << std::endl
69         << "  grid size: " << n << std::endl
70         << "  workers: " << threads << std::endl
71         << "  tolerance: " << tolerance << std::endl;
72 }
73
74 // instantiate a problem
75 Example problem("cliche", 1.0, processors, n);
76
77 // instantiate a solver
78 Jacobi solver(tolerance, threads);
79 // solve
80 solver.solve(problem);
81 // save the results
82 Visualizer vis;
83 vis.csv(problem);
84
85 // initialize mpi
86 status = MPI_Finalize();
87 if (status) {
88     throw("error in MPI_Finalize");
89 }
90
91 // all done
92 return 0;
93 }
```

# The Jacobi declaration

```
1 class acm114::laplace::Jacobi : public acm114::laplace::Solver {
2     // interface
3 public:
4     virtual void solve(Problem &);
5
6     // meta methods
7 public:
8     inline Jacobi(double tolerance, size_t workers);
9     virtual ~Jacobi();
10
11     // data members
12 private:
13     double _tolerance;
14     size_t _workers;
15
16     // disable these
17 private:
18     Jacobi(const Jacobi &);
19     const Jacobi & operator= (const Jacobi &);
20 };
```

# The Problem declaration

```
1 class acml14::laplace::Problem {
2     //typedefs
3 public:
4     typedef std::string string_t;
5     // interface
6 public:
7     string_t name() const;
8     inline MPI_Comm communicator() const;
9     inline int rank() const;
10    // access to my grid
11    inline Grid & solution();
12    inline const Grid & solution() const;
13    // interface used by the solver
14    virtual void initialize();
15    virtual void applyBoundaryConditions() = 0;
16    // meta methods
17 public:
18     Problem(string_t name, double interval, int processors, size_t points);
19     virtual ~Problem();
20    // data members
21 protected:
22     string_t _name;
23     double _delta, _x0, _y0;
24     int _rank, _size, _processors;
25     int _place[2];
26     MPI_Comm _cartesian;
27     Grid _solution;
28    // disable these
29 private:
30     Problem(const Problem &);
31     const Problem & operator= (const Problem &);
32 };
```

# The Problem constructor

```
94 Problem::Problem(  
95     string_t name, double interval, int processors, size_t points) :  
96     _name(name),  
97     _delta(interval/((points-2)*processors+1)),  
98     _x0(0.0), _y0(0.0),  
99     _rank(0), _size(0), _processors(processors), _place(),  
100     _cartesian(),  
101     _solution(points) {  
102  
103     // build the intended layout  
104     int layout[] = { processors, processors };  
105     // find my rank in the world communicator  
106     int worldRank;  
107     MPI_Comm_rank(MPI_COMM_WORLD, &worldRank);  
108     // build a Cartesian communicator  
109     int periods[] = { 0, 0 };  
110     MPI_Cart_create(MPI_COMM_WORLD, 2, &layout[0], periods, 1, &_cartesian);  
111     // check whether i can participate  
112     if (_cartesian != MPI_COMM_NULL) {  
113         // get my rank in the cartesian communicator  
114         MPI_Comm_rank(_cartesian, &_rank);  
115         MPI_Comm_size(_cartesian, &_size);  
116         // get my logical position on the process grid  
117         MPI_Cart_coords(_cartesian, _rank, 2, &_place[0]);  
118         // now compute my offset in physical space  
119         _x0 = 0.0 + (points-2)*_place[0]*_delta;  
120         _y0 = 0.0 + (points-2)*_place[1]*_delta;  
121     } else {  
122         // i was left out because the total number of processors is not a square  
123         std::cout  
124             << "world rank " << worldRank << ": not a member of the cartesian communicator "  
125             << std::endl;  
126     }  
127 }
```



# The Example declaration

```
1 class acml14::laplace::Example : public acml14::laplace::Problem {
2     // interface
3 public:
4     virtual void applyBoundaryConditions();
5
6     // meta methods
7 public:
8     inline Example(
9         string_t name, double interval, int processors, size_t points);
10    virtual ~Example();
11
12    // disable these
13 private:
14     Example(const Example &);
15     const Example & operator= (const Example &);
16 };
```