

ACM/CS 114

Parallel algorithms for scientific applications

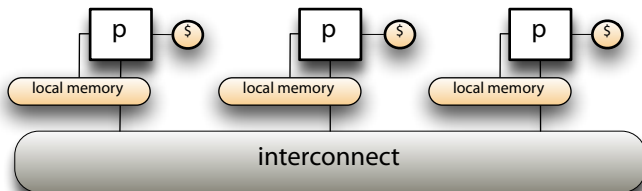
Michael A. G. Aïvázis

California Institute of Technology

Winter 2010

Distributed memory parallelism

- ▶ recall the generic layout of a distributed memory machine



- ▶ each processor has its own private memory space
- ▶ processors communicate via the interconnect substrate
- ▶ the programming model
 - ▶ program consists of a collection of p named processes
 - ▶ each process has its own instruction stream and address space
 - ▶ logically shared data must be partitioned among the processors
 - ▶ communication and synchronization must be orchestrated explicitly
 - ▶ processes communicate via explicit data exchanges

- ▶ the *de facto* standard for writing parallel programs using message passing
 - ▶ a library of routines callable from almost any programming language
 - ▶ that enables communication among multiple processes
 - ▶ standardized and portable API with good implementations available for almost any kind of parallel computer
- ▶ MP I is large and complex
 - ▶ more than 125 functions, lot's of options and communication protocols
 - ▶ but for most practical purposes, a small subset will suffice
 - ▶ short introduction today, more when we consider specific physics
- ▶ two major versions available – check your installation for compliance
 - ▶ MP I-1: parallel machine management, process groups, collective operations, point-to-point operations, virtual topologies, profiling
 - ▶ MP I-2: dynamic process management, one-sided operations, parallel I/O, (simplistic) bindings for C++
- ▶ `openmpi`: currently the best open source implementation
 - ▶ well-architected, thread safe, fast, decent support from a broad community

Getting started

- ▶ **compiling and linking:**
 - ▶ most MPI implementation supply wrappers around the available compilers
 - ▶ e.g. `mpicc`, `mpic++`, `mpif77`, `mpif90`
 - ▶ it's not magic, so you can do it on your own to
 - ▶ override the system defaults (without upsetting the sysadmins...)
 - ▶ build multiple versions so you can benchmark
- ▶ **staging and launching:**
 - ▶ most implementations provide `mpirun` to
 - ▶ control the total number of desired processes
 - ▶ specify the hostnames of the machines to use
 - ▶ specify the mapping of processes to machines/CPU's/cores
 - ▶ establish the current working directory, if possible, for all processes
 - ▶ launch the program
 - ▶ but most installations do not permit its use; they have queuing systems instead
 - ▶ PBS, LSF, torque, maui, ...
 - ▶ specified and documented in the “welcome” package of most supercomputer centers
 - ▶ scheduling of jobs, guarantee exclusive access to your allocated machines, establish upper time limit, charge the right account for your uses

At runtime

- initializing the coöperating processes:

```
1 int MPI_Init(int* argc, char ***argv);
```

- ▶ note the strange signature; see Slide 7 for an example of its use
- ▶ some implementations – notably MPICH, the reference implementation – used command line arguments to pass information from `mpirun` to the runtime environment
- ▶ so they need *write* access to the command line arguments to strip the extras
- ▶ thankfully, not done any more
- ▶ must be the first `MPI` in your program; nothing is initialized correctly until it returns
 - ▶ if this call does not return `MPI_SUCCESS`, you should abort
- ▶ don't forget to shut everything down:

```
1 int MPI_Finalize(void);
```

- ▶ must be the last MPI call in your program; nothing is in usable state after it returns

Groups and communicators

- ▶ every MPI process belongs to at least one *group*
- ▶ groups have associated *communicators* that provide the context for data exchanges and synchronization among processes
- ▶ processes in a given communicator get *ranked*
 - ▶ a communicator of p processes assigns ranks 0 through $p - 1$
 - ▶ a process can discover the communicator size and its own rank by using

```
1 int MPI_Comm_size(MPI_Comm communicator, int* size);  
2 int MPI_Comm_rank(MPI_Comm communicator, int* rank);
```

- ▶ the MPI runtime environment creates the *global* communicator
 - ▶ known as MPI_COMM_WORLD
 - ▶ all processes are members
- ▶ it is good practice to learn to manage your own
 - ▶ to narrow down global operations to processor subsets
 - ▶ to promote *reuse*
 - ▶ more details later

Hello world

```
1 #include <mpi.h>
2 #include <stdio.h>
3
4 int main(int argc, char* argv[]) {
5     int status;
6     int rank, size;
7
8     /* initialize MPI */
9     status = MPI_Init(&argc, &argv);
10    if (status != MPI_SUCCESS) {
11        printf("error in MPI_Init; aborting...\n");
12        return status;
13    }
14
15    /* all good -- get process info and display it */
16    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
17    MPI_Comm_size(MPI_COMM_WORLD, &size);
18    printf("hello from %03d/%03d!\n", rank, size);
19
20    /* shut down MPI */
21    MPI_Finalize();
22
23    return 0;
24 }
```

Messages

- ▶ in general, data exchanges through MPI calls involve
 - ▶ a communicator
 - ▶ specifies which processes participate in the exchange
 - ▶ resolves process ranks into processes
 - ▶ *collective* operations involve the entire communicator
 - ▶ *point-to-point* operations require the rank of the message source or destination
 - ▶ the details of the message payload
 - ▶ the address of the source buffer
 - ▶ the data type of the buffer contents
 - ▶ the number of items in the buffer
- ▶ MPI provides some data abstractions to
 - ▶ hide machine dependencies in the data representations to enhance portability and support heterogeneous clusters
 - ▶ support user defined data types
 - ▶ support non-contiguous data layouts

Collective operations: global reductions

- ▶ *collective* operations involve all processes in a given communicator
- ▶ the MPI version of our global reduction example uses

```
1 int MPI_Allreduce(  
2     void* send_buffer, void* recv_buffer,  
3     int count, MPI_Datatype datatype, MPI_Op operation,  
4     MPI_Comm communicator  
5 );
```

- ▶ example legal values for MPI_Datatype
 - ▶ C: MPI_INT, MPI_LONG, MPI_DOUBLE
 - ▶ FORTRAN: MPI_INTEGER, MPI_DOUBLE_PRECISION, MPI_COMPLEX
- ▶ legal values for MPI_Op
 - ▶ MPI_MAX, MPI_MIN, MPI_MAXLOC, MPI_MINLOC
 - ▶ MPI_SUM, MPI_PROD
 - ▶ MPI_LAND, MPI_LOR, MPI_LXOR
 - ▶ MPI_BAND, MPI_BOR, MPI_BXOR
 - ▶ MPI_REPLACE

Example reduction using MPI

```
1 #include <mpi.h>
2 #include <stdio.h>
3
4 int main(int argc, char* argv[]) {
5     int status;
6     int rank;
7     int square, sum;
8
9     /* initialize MPI */
10    status = MPI_Init(&argc, &argv);
11    if (status != MPI_SUCCESS) {
12        printf("error in MPI_Init; aborting...\n");
13        return status;
14    }
15
16    /* get the process rank */
17    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
18    /* form the square */
19    square = rank*rank;
20    /* each process contributes the square of its rank */
21    MPI_Allreduce(&square, &sum, 1, MPI_INT, MPI_SUM, MPI_COMM_WORLD);
22    /* print out the result */
23    printf("%03d: sum = %d\n", rank, sum);
24
25    /* shut down MPI */
26    MPI_Finalize();
27
28    return 0;
29 }
```

Point to point communication

- ▶ to send a message

```
1 int MPI_Send(  
2     void* buffer, int count, MPI_Datatype datatype,  
3     int destination, int tag, MPI_Comm communicator  
4 );
```

- ▶ to receive a message

```
1 int MPI_Recv(  
2     void* buffer, int count, MPI_Datatype datatype,  
3     int source, int tag, MPI_Comm communicator  
4 );
```

- ▶ the tag enables choosing the order you may receive pending messages
- ▶ but for a given (source,tag,communicator) messages are received in the order they were sent
- ▶ receiving via wildcards: MPI_ANY_SOURCE and MPI_ANY_TAG
- ▶ in *standard* communication mode, sending and receiving messages are *blocking*, so the function does not return until you can safely access the buffer
 - ▶ to read, free, etc.

Communication modes

- ▶ in standard mode, the specification does not explicitly mention buffering strategy
 - ▶ buffering messages would remove some of the access constraints but it requires time and storage for the multiple copies
 - ▶ portability across implementations implies conservative assumptions about the order of initiation of sends and receives to avoid deadlock
- ▶ in *ready* mode, you must post a receive before the matching send can be initiated
 - ▶ `MPI_Rsend`, `MPI_Rrecv`
- ▶ in *buffered* mode, sends can be initiated, and may complete, regardless of when the matching receive is initiate
 - ▶ `MPI_Bsend`, `MPI_Brecv`
- ▶ in *synchronous* mode, sends can be initiated regardless of whether the matching receive has been initiated, but the send will not return until the message has been received
 - ▶ `MPI_Ssend`, `MPI_Srecv`

Asynchronous communication

- ▶ there are non-blocking versions of all these

```
1 int MPI_Isend(  
2     void* buffer, int count, MPI_Datatype datatype,  
3     int destination, int tag,  
4     MPI_Comm communicator, MPI_Request* request  
5 );
```

- ▶ faster, but you must take care to not access the message buffers until the messages have been delivered
 - ▶ more details later in the course, as needed
- ▶ for sends
 - ▶ standard mode: `MPI_Isend`
 - ▶ ready mode: `MPI_Irsend`
 - ▶ buffered mode: `MPI_Ibsend`
 - ▶ synchronous mode: `MPI_Issend`
- ▶ only one call for receives: `MPI_Irecv`
- ▶ extra `request` argument to check for completion of the request
 - ▶ `MPI_Test`, `MPI_Wait` and their relatives

Creating communicators and groups

- ▶ communicators and groups are intertwined
 - ▶ you cannot create a group without a communicator
 - ▶ you cannot create a communicator without a group
- ▶ the cycle is broken by `MP I_COMM_WORLD`

```
1 #include <mpi.h>
2
3 int main(int argc, char* argv[]) {
4     /* declare a communicator and a couple of groups */
5     MPI_Comm workers;
6     MPI_Group world_grp, workers_grp;
7
8     /* initialize MPI; for brevity all status checks are omitted */
9     MPI_Init(&argc, &argv);
10
11     /* get the world communicator to build its group */
12     MPI_Comm_group(MPI_COMM_WORLD, &world_grp);
13
14     /* build another group by excluding a process */
15     MPI_Group_excl(world_grp, 1, 0, &workers_grp);
16
17     /* now build a communicator out of the processes in workers_grp */
18     MPI_Comm_create(MPI_COMM_WORLD, worker_grp, &workers);
19
20     /* etc... */
21
22     /* shut down MPI */
23     MPI_Finalize();
24
25     return 0;
26 }
```

Manipulating communicators and groups

- ▶ releasing resources

```
1 int MPI_Group_free(MPI_Group* group);  
2 int MPI_Comm_free(MPI_Comm* communicator);  
3 int MPI_Comm_disconnect(MPI_Comm* communicator);
```

- ▶ you can make a new group by adding or removing processes from an existing one

```
1 int MPI_Group_incl(  
2     MPI_Group grp, int n, int* ranks, MPI_Group* new_group);  
3 int MPI_Group_excl(  
4     MPI_Group grp, int n, int* ranks, MPI_Group* new_group);
```

- ▶ or by using set operations

```
1 int MPI_Group_union(  
2     MPI_Group grp1, MPI_Group grp2, MPI_Group* new_group);  
3 int MPI_Group_intersection(  
4     MPI_Group grp1, MPI_Group grp2, MPI_Group* new_group);  
5 int MPI_Group_difference(  
6     MPI_Group grp1, MPI_Group grp2, MPI_Group* new_group);
```

- ▶ the function

```
1 double MPI_Wtime();
```

returns the time in seconds from some arbitrary time in the past

- ▶ guaranteed not to change only for the duration of the process
- ▶ you can compute the elapsed time for any program segment by making calls at the beginning and the end and computing the difference
- ▶ no guarantees about synchronized clocks among different processes
- ▶ you can compute the clock resolution by using

```
1 double MPI_Wtick();
```


Other collective operations

- ▶ `MPI_Scan` computes partial reductions: the p^{th} process receives the result from processes 0 through $p - 1$

```
1 int MPI_Scan(  
2     void* send_buffer, void* recv_buffer,  
3     int count, MPI_Datatype datatype, MPI_Op operation,  
4     MPI_Comm communicator  
5 );
```

- ▶ `MPI_Reduce` collects the result at only the given process root

```
1 int MPI_Reduce(  
2     void* send_buffer, void* recv_buffer,  
3     int count, MPI_Datatype datatype, MPI_Op operation,  
4     int root, MPI_Comm communicator  
5 );
```

- ▶ synchronization is also a global operation:


```
1 int MPI_Barrier(MPI_Comm communicator);
```

participating processes block at a barrier until they have all reached it

Scatter

- ▶ `MPI_Scatter` sends data from `root` to all processes

```
1 int MPI_Scatter(  
2     void* send_buffer, int send_count, MPI_Datatype send_datatype,  
3     void* recv_buffer, int recv_count, MPI_Datatype recv_datatype,  
4     int root, MPI_Comm communicator  
5 );
```




- ▶ it is as if the data in `send_buffer` were split in p segments, and the i^{th} process receives the i^{th} segment
- ▶ the `send_xxx` arguments are only meaningful for `root`; they are ignored for other processes
- ▶ the arguments `root` and `communicator` must be passed identical values by all processes

Gather

- ▶ the converse is `MPI_Gather` with `root` receiving data from all processes

```
1 int MPI_Gather(  
2     void* send_buffer, int send_count, MPI_Datatype send_datatype,  
3     void* recv_buffer, int recv_count, MPI_Datatype recv_datatype,  
4     int root, MPI_Comm communicator  
5 );
```



- ▶ it is as if p messages, one from each processes, were concatenated in rank order and placed at `recv_buffer`
- ▶ the `recv_xxx` arguments are only meaningful for `root`; they are ignored for other processes
- ▶ the arguments `root` and `communicator` must be passed identical values by all processes

Broadcasting operations


- ▶ `MPI_Alltoall` sends data from all processes to all processes in a global scatter/gather

```
1 int MPI_Alltoall(  
2     void* send_buffer, int send_count, MPI_Datatype send_datatype,  
3     void* recv_buffer, int recv_count, MPI_Datatype recv_datatype,  
4     MPI_Comm communicator  
5 );
```

- ▶ use `MPI_Bcast` to send the contents of a buffer from `root` to all processes in a communicator

```
1 int MPI_Bcast(  
2     void* buffer, int count, MPI_Datatype datatype,  
3     int root, MPI_Comm communicator  
4 );
```

Data movement patterns for the collective operations



Virtual topologies

