

# ACM/CS 114

## Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Winter 2010

# Printing out the initial grid

- ▶ we should be able to print out the initialized grid

```
1 #> mm laplace
2 #> laplace
3 #> cat laplace.csv
4 0,0,0.3827,0.7071,0.9239,1,0.9239,0.7071,0.3827,1.225e-16
5 1,0,1,1,1,1,1,1,1,0
6 2,0,1,1,1,1,1,1,1,0
7 3,0,1,1,1,1,1,1,1,0
8 4,0,1,1,1,1,1,1,1,0
9 5,0,1,1,1,1,1,1,1,0
10 6,0,1,1,1,1,1,1,1,0
11 7,0,1,1,1,1,1,1,1,0
12 8,0,0.01654,0.0306,0.03992,0.04321,0.0399,0.03056,0.01654,0
```

- ▶ notice that
  - ▶ the top line contains some recognizable values
  - ▶ the left and right borders are set to zero
  - ▶ the interior of the grid is painted with our initial guess
- ▶ still to do:
  - ▶ write the update
  - ▶ build a grid with the exact solution
  - ▶ build the error field (why?)

# Fleshing out the solver

```
169 // the solver driver
170 void laplace(Grid & current, double tolerance) {
171     // create and initialize temporary storage
172     Grid next(current.size());
173     initialize(next);
174     // put an upper bound on the number of iterations
175     long max_iterations = (long) 1e4;;
176     for (long iterations = 0; iterations < max_iterations; iterations++) {
177         double max_dev = 0.0;
178         // do an iteration step
179         // leave the boundary alone
180         // iterate over the interior of the grid
181         for (size_t j=1; j < current.size()-1; j++) {
182             for (size_t i=1; i < current.size()-1; i++) {
183                 // update
184                 next(i,j) = 0.25*(
185                     current(i+1,j)+current(i-1,j)+current(i,j+1)+current(i,j-1));
186                 // compute the deviation from the last generation
187                 double dev = std::abs(next(i,j) - current(i,j));
188                 // and update the maximum deviation
189                 if (dev > max_dev) {
190                     max_dev = dev;
191                 }
192             }
193         }
194         // swap the blocks between the two grids
195         Grid::swapBlocks(current, next);
196         // check convergence
197         if (max_dev < tolerance) {
198             break;
199         }
200     }
201     return;
202 }
```

# Adding the new grid interface

- here is the declaration of `Grid::swapBlocks`

```
30 class Grid {  
31     // interface  
32     public:  
33     ...  
34     // exchange the data blocks of two compatible grids  
35     static void swapBlocks(Grid &, Grid &);  
36     ...  
37 };
```

- and its definition

```
69 void Grid::swapBlocks(Grid & g1, Grid & g2) {  
70     // bail out if the two operands are not compatible  
71     if (g1.size() != g2.size()) {  
72         throw "Grid::swapblocks: size mismatch";  
73     }  
74     if (g1.delta() != g2.delta()) {  
75         throw "Grid::swapblocks: spacing mismatch";  
76     }  
77     // but if they are, just exchange their data buffers  
78     double * temp = g1._block;  
79     g1._block = g2._block;  
80     g2._block = temp;  
81     // all done  
82     return;  
83 }
```

# Reworking the driver

```
239 // build a visualizer
240 Visualizer vis;
241
242 // compute the exact solution
243 Grid solution(N);
244 exact(solution);
245 std::fstream exact_stream("exact.csv", std::ios_base::out);
246 vis.csv(solution, exact_stream);
247
248 // allocate space for the solution
249 Grid potential(N);
250 // initialize and apply our boundary conditions
251 initialize(potential);
252 // call the solver
253 laplace(potential, tolerance);
254 // open a stream to hold the answer
255 std::fstream output_stream(filename, std::ios_base::out);
256 // build a visualizer and render the solution in our chosen format
257 vis.csv(potential, output_stream);
258
259 // compute the error field
260 Grid error(N);
261 relative_error(potential, solution, error);
262 std::fstream error_stream("error.csv", std::ios_base::out);
263 vis.csv(error, error_stream);
264
265 // all done
266 return 0;
267 }
```

# Computing the exact solution and the error field

```
143 void exact(Grid & grid) {
144     // paint the exact solution
145     for (size_t j=0; j < grid.size(); j++) {
146         for (size_t i=0; i < grid.size(); i++) {
147             double x = i*grid.delta();
148             double y = j*grid.delta();
149             grid(i,j) = std::exp(-pi*y)*std::sin(pi*x);
150         }
151     }
152     return;
153 }
154
155 void relative_error(
156     const Grid & computed, const Grid & exact, Grid & error) {
157     // compute the relative error
158     for (size_t j=0; j < exact.size(); j++) {
159         for (size_t i=0; i < exact.size(); i++) {
160             if (exact(i,j) == 0.0) { // hm... sloppy!
161                 error(i,j) = std::abs(computed(i,j));
162             } else {
163                 error(i,j) = std::abs(computed(i,j) - exact(i,j))/exact(i,j);
164             }
165         }
166     }
167     return;
168 }
```

# Shortcomings

- ▶ numerics:
  - ▶ it converges very slowly; other update *schemes* improve on this
  - ▶ our approximation is very low order, so it takes very large grids to produce a few digits of accuracy
  - ▶ the convergence criterion has some unwanted properties; it triggers
    - ▶ prematurely: large swaths of constant values may never get updated
    - ▶ it would trigger even if we were updating the wrong grid!
- ▶ design:
  - ▶ separate the problem specification from its solution
  - ▶ there are other objects lurking, waiting to be uncovered
  - ▶ someone should make the graphic visualizer
  - ▶ restarts anybody?
  - ▶ how would you try out different convergence criteria? update schemes? memory layouts?
- ▶ usability:
  - ▶ supporting interchangeable parts requires damage to the top level driver
    - ▶ to enable the user to make the selection
    - ▶ to expose new command line arguments that configure the new parts

# Parallelization using threads

- ▶ the shared memory implementation requires
  - ▶ a scheme so that threads can update cells without the need for locks
  - ▶ while maximizing locality of data access
  - ▶ even the computation of the convergence criterion can be parallelized
- ▶ parallelization strategy
  - ▶ we will focus on parallelizing the iterative grid update
    - ▶ grid initialization, visualization, computing the exact answer and the error field do not depend on the *number of iterations*
  - ▶ the finest grain of work is clearly an individual cell update based on the value of its four nearest neighbors
  - ▶ for this two dimensional example, we can build coarser grain tasks using
    - ▶ horizontal or vertical strips
    - ▶ non-overlapping blocks
    - ▶ the strategy gets more complicated if you want to perform the update in place
  - ▶ the communication patterns are trivial for the double buffering layout; only the final update of the convergence criterion requires any locking
  - ▶ each coarse grain task can be assigned to a thread



# Required changes to the sequential solution

- ▶ what is needed
  - ▶ an object to hold the problem information shared among the threads
  - ▶ the per-thread administrative data structure that holds the thread id and the pointer to the shared information
    - ▶ this is the argument to `pthread_create`
  - ▶ a mutex to protect the update of the global convergence criterion
  - ▶ a `pthread_create` compatible worker routine
  - ▶ a change at the top-level driver to enable the user to choose the number of threads
- ▶ and a strategy for managing the thread life cycle
  - ▶ synchronization is trivial if
    - ▶ we spawn our threads to perform the updates of a single iteration
    - ▶ harvest them
    - ▶ check the convergence criterion
    - ▶ stop, or respawn them if another iteration is necessary
  - ▶ can the convergence test be done in parallel?
    - ▶ so we don't have to pay the create/harvest overhead?