

ACM/CS 114

Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Winter 2010

Dense matrix problems

- ▶ we'll take a look at
 - ▶ inner and outer products of two vector
 - ▶ matrix-vector and matrix-matrix multiplication
 - ▶ LU and Cholesky decompositions
 - ▶ QR factorizations
 - ▶ computing eigenvalues and eigenvectors
 - ▶ fast Fourier transforms
- ▶ when solving a problem of size n on p processors, we will assume
 - ▶ that p , and occasionally \sqrt{p} divides n
 - ▶ that p is a perfect square, when forming two-dimensional process grids
 - ▶ matrices are $n \times n$ – square, not rectangular
 - ▶ we are memory constrained and data replication must be minimized
- ▶ these problems have been studied extensively and form the core of scientific computing on parallel machines
 - ▶ excellent implementations available
 - ▶ interest has been revived due to the expected disruption by multi-core architectures

Vector inner product

- ▶ the inner product of two n -vectors x, y is given by

$$x^T y = \sum_{i=1}^n x_i y_i$$

which requires n multiplications and $n - 1$ additions

- ▶ parallelization strategy:
 - ▶ n fine grain tasks, numbered $i = 1, \dots, n$, that store x_i and y_i , and compute $x_i y_i$
 - ▶ communication is a sum reduction over n fine grain tasks
 - ▶ coarsening is achieved by coalescing n/p tasks together, assuming that each process can accommodate the data storage requirements
 - ▶ and mapping each coarse grain task to a process

Vector outer product

- ▶ the outer product of two n -vectors x and y is the $n \times n$ matrix A given by

$$A_{ij} = x_i y_j$$

which requires n^2 multiplications

- ▶ parallelization strategies are determined by the storage requirements
 - ▶ build a two-dimensional grid of n^2 fine grain tasks numbered (i, j) , with $i, j = 1, \dots, n$; each one computes $x_i y_j$
 - ▶ assuming no data replication is allowed
 - ▶ let task $(i, 1)$ store x_i and task $(1, i)$ store y_i
 - ▶ or, let task (i, i) store both x_i and y_i
 - ▶ either way, the task that owns each element must broadcast it to the other tasks: x_i along the i^{th} task row, y_j along the j^{th} task column
 - ▶ coarsening to p tasks can be accomplished by
 - ▶ combining n/p rows or columns
 - ▶ forming $(n/\sqrt{p}) \times (n/\sqrt{p})$ grid of fine grain tasks
 - ▶ and each coarse grain task can be assigned to a process
- ▶ either way, naïve broadcasting of the components of x and y would require as much total memory as replication
 - ▶ storage can be reduced by circulating portions of x and y through the tasks, with each task using the available portion and passing it on

The product of a matrix with a vector

- ▶ given an $n \times n$ matrix A and an n -vector x , the matrix vector product yields an n -vector y whose components are given by

$$y_i = \sum_{j=1}^n A_{ij}x_j$$

requiring a total of n^2 multiply-add operations

- ▶ once again, the parallelization strategy is determined by how the data is distributed among fine grain tasks
 - ▶ build a two-dimensional grid of n^2 fine grain tasks numbered (i,j) , with $i, j = 1, \dots, n$; each one computes $A_{ij}x_j$
 - ▶ task (i,j) has $a_{i,j}$, but if no data replication is allowed
 - ▶ let task $(i, 1)$ store x_i and task $(1, i)$ store y_i
 - ▶ or, let task (i, i) store both x_i and y_i
 - ▶ the task that owns x_j must broadcast it along the j^{th} task row, and y_i is formed by sum reduction along the i^{th} task column
 - ▶ coarsening into p tasks can be accomplished by combining n/p rows/columns, or by forming $(n/\sqrt{p}) \times (n/\sqrt{p})$ blocks
 - ▶ and each coarse grain task can be assigned to a process

Coarsening along rows or columns

- ▶ for one-dimensional coarsening into n/p task rows
 - ▶ if x is stored in one task, it must be broadcast to all others
 - ▶ if x is distributed among tasks, with n/p components per task, then multiple broadcasts are required
 - ▶ each task computes the inner product of its n/p rows of A with the entire x to produce n/p components of y
- ▶ for one-dimensional coarsening into n/p task columns
 - ▶ n/p components of x are distributed among the tasks
 - ▶ each task computes the linear combination of its n/p columns with coefficients from its copy of x
 - ▶ since the right parts of x are already available, no communication is required
 - ▶ y is generated by a sum reduction across tasks
- ▶ these two are *duals* of each other
 - ▶ row coarsening begins with broadcast, followed by communication-free inner products
 - ▶ column coarsening begins with communication-free linear combinations, follows by a reduction

Two dimensional coarsening

- ▶ for two dimensional coarsening, we form $(n/\sqrt{p}) \times (n/\sqrt{p})$ blocks of fine grain task
 - ▶ each one holding a $(n/\sqrt{p}) \times (n/\sqrt{p})$ block of A
 - ▶ with components of x distributed either across one task row, or along the diagonal, n/p components per task

S