

# ACM/CS 114

## Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Winter 2010

# Dense matrix problems

- ▶ we'll take a look at
  - ▶ inner and outer products of two vector
  - ▶ matrix-vector and matrix-matrix multiplication
  - ▶  $LU$  factorization and Cholesky decomposition
  - ▶ QR factorization
  - ▶ computing eigenvalues and eigenvectors
  - ▶ fast Fourier transforms
- ▶ when solving a problem of size  $n$  on  $p$  processors, we will assume
  - ▶ that  $p$ , and occasionally  $\sqrt{p}$  divides  $n$
  - ▶ that  $p$  is a perfect square, when forming two-dimensional process grids
  - ▶ matrices are  $n \times n$  – square, not rectangular
  - ▶ we are memory constrained and data replication must be minimized
- ▶ these problems have been studied extensively and form the core of scientific computing on parallel machines
  - ▶ excellent implementations available
  - ▶ interest has been revived due to the expected disruption by multi-core architectures

# Vector inner product

- ▶ the inner product of two  $n$ -vectors  $x, y$  is given by

$$x^T y = \sum_{i=1}^n x_i y_i \quad (1)$$

which requires  $n$  multiplications and  $n - 1$  additions

- ▶ parallelization strategy:
  - ▶  $n$  fine grain tasks, numbered  $i = 1, \dots, n$ , that store  $x_i$  and  $y_i$ , and compute  $x_i y_i$
  - ▶ communication is a sum reduction over  $n$  fine grain tasks
  - ▶ coarsening is achieved by coalescing  $n/p$  tasks together, assuming that each process can accommodate the data storage requirements
  - ▶ and mapping each coarse grain task to a process

# Vector outer product

- ▶ the outer product of two  $n$ -vectors  $x$  and  $y$  is the  $n \times n$  matrix  $A$  given by

$$A_{ij} = x_i y_j \quad (2)$$

which requires  $n^2$  multiplications

- ▶ parallelization strategies are determined by the storage requirements
  - ▶ build a two-dimensional grid of  $n^2$  fine grain tasks numbered  $(i, j)$ , with  $i, j = 1, \dots, n$ ; each one computes  $x_i y_j$
  - ▶ assuming no data replication is allowed
    - ▶ let task  $(i, 1)$  store  $x_i$  and task  $(1, i)$  store  $y_i$
    - ▶ or, let task  $(i, i)$  store both  $x_i$  and  $y_i$
  - ▶ either way, the task that owns each element must broadcast it to the other tasks:  $x_i$  along the  $i^{\text{th}}$  task row,  $y_j$  along the  $j^{\text{th}}$  task column
  - ▶ coarsening to  $p$  tasks can be accomplished by
    - ▶ combining  $n/p$  rows or columns
    - ▶ forming  $(n/\sqrt{p}) \times (n/\sqrt{p})$  grid of fine grain tasks
  - ▶ and each coarse grain task can be assigned to a process
- ▶ either way, naïve broadcasting of the components of  $x$  and  $y$  would require as much total memory as replication
  - ▶ storage can be reduced by circulating portions of  $x$  and  $y$  through the tasks, with each task using the available portion and passing it on

# The product of a matrix with a vector

- ▶ given an  $n \times n$  matrix  $A$  and an  $n$ -vector  $x$ , the matrix vector product yields an  $n$ -vector  $y$  whose components are given by

$$y_i = \sum_{j=1}^n A_{ij}x_j \quad (3)$$

requiring a total of  $n^2$  multiply-add operations

- ▶ once again, the parallelization strategy is determined by how the data is distributed among fine grain tasks
  - ▶ build a two-dimensional grid of  $n^2$  fine grain tasks numbered  $(i,j)$ , with  $i, j = 1, \dots, n$ ; each one computes  $A_{ij}x_j$
  - ▶ task  $(i,j)$  has  $A_{i,j}$ , but if no data replication is allowed
    - ▶ let task  $(i, 1)$  store  $x_i$  and task  $(1, i)$  store  $y_i$
    - ▶ or, let task  $(i, i)$  store both  $x_i$  and  $y_i$
  - ▶ the task that owns  $x_j$  must broadcast it along the  $j^{\text{th}}$  task row, and  $y_i$  is formed by sum reduction along the  $i^{\text{th}}$  task column
  - ▶ coarsening into  $p$  tasks can be accomplished by combining  $n/p$  rows/columns, or by forming  $(n/\sqrt{p}) \times (n/\sqrt{p})$  blocks
  - ▶ and each coarse grain task can be assigned to a process

# Coarsening along rows or columns

- ▶ for one-dimensional coarsening into  $n/p$  task rows
  - ▶ if  $x$  is stored in one task, it must be broadcast to all others
  - ▶ if  $x$  is distributed among tasks, with  $n/p$  components per task, then multiple broadcasts are required
  - ▶ each task computes the inner product of its  $n/p$  rows of  $A$  with the entire  $x$  to produce  $n/p$  components of  $y$
- ▶ for one-dimensional coarsening into  $n/p$  task columns
  - ▶  $n/p$  components of  $x$  are distributed among the tasks
  - ▶ each task computes the linear combination of its  $n/p$  columns with coefficients from its copy of  $x$
  - ▶ since the right parts of  $x$  are already available, no communication is required
  - ▶  $y$  is generated by a sum reduction across tasks
- ▶ these two are *duals* of each other
  - ▶ row coarsening begins with broadcast, followed by communication-free inner products
  - ▶ column coarsening begins with communication-free linear combinations, follows by a reduction

# Two dimensional coarsening

- ▶ for two dimensional coarsening, we form  $(n/\sqrt{p}) \times (n/\sqrt{p})$  blocks of fine grain task
  - ▶ each one holding a  $(n/\sqrt{p}) \times (n/\sqrt{p})$  block of  $A$
  - ▶ with components of  $x$  distributed either across one task row, or along the diagonal,  $n/p$  components per task
- ▶ the algorithm combines the features of row/column coarsening
  - ▶ components of  $x$  are broadcast along task columns
  - ▶ each task performs  $n^2/p$  multiplications locally and sums  $n/\sqrt{p}$  sets of products
  - ▶ sum reductions along task rows produce the components of  $y$  by combining the component products

# Matrix multiplication

- ▶ the product of two  $n \times n$  matrices  $A$  and  $B$  is an  $n \times n$  matrix  $C$  given by

$$C_{ij} = \sum_{k=1}^n A_{ik} B_{kj} \quad (4)$$

where each one of  $n^2$  entries requires  $n$  multiply-adds for a total of  $n^3$  operations

- ▶ matrix multiplication can be viewed as
  - ▶  $n^2$  inner products
  - ▶ the sum of  $n$  outer products
  - ▶  $n$  matrix vector products
- ▶ each one produces a parallel algorithm for matrix multiplication
- ▶ but we'll explore a direct solution instead



# Partitioning and communication patterns

- ▶ we build a three dimensional array of  $n^3$  fine grain tasks
  - ▶ with  $i, j, k = 1, \dots, n$ , let task  $(i, j, k)$  be responsible for computing the product  $A_{ij}B_{jk}$
  - ▶ assuming no data replication, we have to distribute the data for  $A$  and  $B$  among  $2n^2$  tasks
  - ▶ suppose that task  $(i, j, j)$  holds  $A_{i,j}$  and task  $(i, j, i)$  holds  $B_{i,j}$
  - ▶ we will refer to tasks along  $i$  and  $j$  as task rows and columns
  - ▶ and tasks along  $k$  as *layers*
- ▶ the communication requirements among tasks are satisfied if we
  - ▶ broadcast the entries of the  $k^{\text{th}}$  column of  $A$  from task  $(i, j, j)$  to each task row in the  $k^{\text{th}}$  layer
  - ▶ broadcast the entries of the  $k^{\text{th}}$  row of  $B$  from task  $(i, j, i)$  to each task column of the  $k^{\text{th}}$  layer
  - ▶ form the result  $C_{ij}$  by the sum reduction of the values held by all the tasks layers  $k$

# Coarsening

- ▶ there are four natural ways to coarsen our  $n \times n \times n$  fine grain tasks into  $p$  coarse grain tasks
  - ▶ by task rows: combine the  $(n/p) \times n \times n$  tasks along a given task row
  - ▶ by task columns: combine the  $n \times (n/p) \times n$  tasks along a given task column
  - ▶ partition the layers in a two dimensional grid by combining  $(n/\sqrt{p}) \times (n/\sqrt{p}) \times n$  fine grain tasks
  - ▶ using three dimensional blocks by combining  $(n/\sqrt[3]{p}) \times (n/\sqrt[3]{p}) \times (n/\sqrt[3]{p})$  tasks
- ▶ the two one dimensional coarsening strategies are similar
  - ▶ for row coarsening
    - ▶ each task needs only the part of  $A$  it already has, but needs all  $B$  entries
    - ▶ so global communication is required to broadcast the  $n^2/p$  entries of  $B$  held by each task
  - ▶ conversely, for column coarsening
    - ▶ each task needs only the parts of  $B$  that it already has, but it needs all of  $A$
    - ▶ so global communication is required to broadcast the  $n^2/p$  entries of  $A$  held by each task
  - ▶ if accumulating  $A$  or  $B$  on each processor is not feasible, tasks can circulate portions of the array in a ring

## Coarsening using a two dimensional grid

- ▶ block matrix multiplication has the same overall form as actual product, with scalar operations replaced by the matrix product of blocks!
- ▶ you should verify that

$$C_{ij} = \sum_{k=1}^{\sqrt{p}} A_{ik} B_{kj} \quad (5)$$

for  $i, j = 1, \dots, \sqrt{p}$

- ▶ assume that task  $(i,j)$  has local access to block  $A_{ij}$  and  $B_{ij}$  and computes block  $C_{ij}$  of the result
- ▶ this requires all blocks  $A_{ik}$  and  $B_{kj}$  for  $k = 1, \dots, \sqrt{p}$  to be communicated
  - ▶ first, a global broadcast of  $A$  blocks across each task row
  - ▶ followed by a global broadcast of  $B$  blocks across each task column
- ▶ memory requirements can be addressed by either of the following:
  - ▶ broadcast blocks of  $A$  across rows while circulating blocks of  $B$  across columns in lock step, so that they arrive at a given task at the same time
  - ▶ circulate blocks of  $A$  horizontally and blocks of  $B$  vertically, after an initial circular shift, so that blocks meet at a given task at the right time

# LU factorization

- ▶ systems of linear equations are ubiquitous in numerical analysis
- ▶ let  $A$  be an  $n \times n$  matrix,  $b$  a known  $n$ -vector; we are looking for  $x$  such that

$$Ax = b \quad (6)$$

- ▶ a commonly used direct method for solving this system is to convert  $A$  into the product of a lower triangular matrix  $L$  with an upper triangular matrix  $U$

$$A = LU \quad (7)$$

known as  $LU$  factorization

- ▶ Eq. 6 becomes

$$LUx = b \quad (8)$$

which we can now solve in two simpler steps

$$Ly = b \quad (9)$$

$$Ux = y \quad (10)$$

where we first solve the lower triangular system by forward substitution, followed by solving the upper triangular system by back substitution to obtain  $x$

# LU by Gaussian elimination

- ▶ we can compute the  $LU$  factorization of  $A$  using Gaussian elimination

---

**Algorithm 1:** LU( $A$ )

---

```
1 for  $k = 1$  to  $n - 1$  do  
2   for  $i = k + 1$  to  $n$  do  
3      $L_{ik} = A_{ik} / A_{kk}$   
4   for  $j = k + 1$  to  $n$  do  
5     for  $i = k + 1$  to  $n$  do  
6        $A_{ij} = A_{ij} - L_{ik}A_{kj}$ 
```

---

which encodes  $L$  and  $U$  in place by overwriting  $A$

- ▶ Alg. 1 requires roughly  $n^3/3$  multiply-adds and  $n^2/2$  divisions
- ▶ we may also need *pivoting* to ensure numerical stability (and existence)
- ▶ Alg. 1 is one of many algorithms expressed essentially as a triply nested loop
  - ▶ the three indices can be ordered in any of  $3!$  ways, with totally different memory access patterns
  - ▶ in parallel, the  $kji$  and  $kji$  forms may be the most efficient

# Parallel $LU$ decomposition

- ▶ number fine grain tasks as  $(i,j)$  with  $i,j = 1, \dots, n$ ; each task
  - ▶ stores  $A_{ij}$
  - ▶ computes and stores  $U_{ij}$ , if  $i \leq j$
  - ▶ computes and stores  $L_{ij}$ , if  $i > j$

yielding a two dimensional array of  $n^2$  tasks

- ▶ no need to compute and store
  - ▶ the zeroes in the lower triangle of  $U$
  - ▶ the unit diagonal and the zeroes in the upper triangle of  $L$
- ▶ in order to create  $p$  coarse grain tasks we could combine
  - ▶  $n/p$  rows or columns of fine grain tasks
  - ▶  $(n/\sqrt{p}) \times (n/\sqrt{p})$  blocks of tasks

and map each one to a process

# Communication patterns for parallel $LU$ decomposition

---

**Algorithm 2:**  $LU(A, \text{task}=(i,j))$ 

---

```
1 for  $k = 1$  to  $\min(i,j) - 1$  do
2   recv  $A_{kj}$ 
3   recv  $L_{ik}$ 
4    $A_{ij} = A_{ij} - L_{ik}A_{kj}$ 
5 if  $i \leq j$  then
6   broadcast  $A_{ij}$  to  $(k,j), k = i + 1, \dots, n$ 
7 else
8   recv  $A_{jj}$ 
9    $L_{ij} = A_{ij}/A_{jj}$ 
10  broadcast  $L_{ij}$  to  $(i,k), k = i + 1, \dots, n$ 
```

---

## Row coarsening

- ▶ with one dimensional row coarsening
  - ▶ we forgo parallelism in updating rows
  - ▶ there is no need to broadcast the multipliers  $L_{ij}$  since each row is contained entirely within a task
  - ▶ we still need the vertical broadcasts of matrix rows to the tasks below

---

**Algorithm 3:** LU( $A$ , task= $(i, j)$ ) by rows

```

1 for  $k = 1$  to  $n - 1$  do
2   if  $k \in \text{myrows}$  then
3     broadcast  $\{A_{kj} : k \leq j \leq n\}$ 
4   else
5     recv  $\{A_{kj} : k \leq j \leq n\}$ 
6   for  $i \in \text{myrows}, i > k$  do
7      $L_{ik} = A_{ik} / A_{kk}$ 
8   for  $j = k + 1$  to  $n$  do
9     for  $i \in \text{myrows}, i > k$  do
10       $A_{ij} = A_{ij} - L_{ik}A_{kj}$ 

```



# Observations on row coarsening

- ▶ each task becomes idle as soon as it last row is completed
  - ▶ if rows are contiguous, a task may finish long before the overall computation is done
  - ▶ even worse, updating rows requires progressively less work with increasing row number
- ▶ we may improve concurrency and load balance
  - ▶ by assigning rows to tasks in a cyclic manner where row  $i$  is updated by task  $i \bmod p$
  - ▶ other mappings may be useful
- ▶ other improvements involve overlapping computation with communication
  - ▶ at step  $k$ , each task completes updating its portion of the remaining unreduced matrix before moving on to step  $k + 1$
  - ▶ however, the task that owns the  $k + 1$  row could broadcast it as soon as it becomes available, before moving on to the step  $k$  update
  - ▶ this *send ahead* strategy may grant other tasks earlier access to the data necessary to start working on the next step

# Column coarsening

---

**Algorithm 4:**  $\text{LU}(A, \text{task}=(i,j))$  by columns

---

```
1 for  $k = 1$  to  $n - 1$  do
2   if  $k \in \text{mycolumns}$  then
3     for  $i = k + 1$  to  $n$  do
4        $L_{ik} = A_{ik} / A_{kk}$ 
5       broadcast  $\{L_{ik} : k < i \leq n\}$ 
6   else
7     recv  $\{L_{ik} : k < i \leq n\}$ 
8   for  $i \in \text{mycolumns}, j > k$  do
9     for  $i = k + 1$  to  $n$  do
10       $A_{ij} = A_{ij} - L_{ik}A_{kj}$ 
```

---

- observations similar to row coarsening apply

# Block coarsening

---

**Algorithm 5:** LU( $A$ , task= $(i, j)$ ) by blocks

---

```
1 for  $k = 1$  to  $n - 1$  do
2   if  $k \in \text{myrows}$  then
3     broadcast  $\{A_{kj} : j \in \text{mycolumns}, j > k\}$  to all tasks in my task
        column
4   else
5     recv  $\{A_{kj} : j \in \text{mycolumns}, j > k\}$ 
6   if  $k \in \text{mycolumns}$  then
7     for  $i \in \text{myrows}, i > k$  do
8        $L_{ik} = A_{ik} / A_{kk}$ 
9     broadcast  $\{L_{ik} : i \in \text{myrows}, i > k\}$  to all tasks in my task row
10  else
11    recv  $\{L_{ik} : i \in \text{myrows}, i > k\}$ 
12  for  $j \in \text{mycolumns}, j > k$  do
13    for  $i \in \text{myrows}, i > k$  do
14       $A + ij = A_{ij} - L_{ik}A_{kj}$ 
```

# Observations on block coarsening

- ▶ each task becomes idle as soon as its last row and column are completed
  - ▶ if rows and columns are in contiguous blocks, a task may finish long before the overall computation is done
  - ▶ even worse, computing multipliers and updating blocks requires progressively less work with increasing row and column numbers
- ▶ we may improve concurrency and load balance
  - ▶ by assigning rows and columns to tasks in a cyclic manner where  $A_{ij}$  is assigned to task  $(i \bmod \sqrt{p}, j \bmod \sqrt{p})$
  - ▶ other mappings may be useful
- ▶ other improvements involve overlapping computation with communication
  - ▶ at step  $k$ , each task completes updating its portion of the remaining unreduced submatrix before moving on to step  $k + 1$
  - ▶ the broadcast of each segment of row  $k + 1$ , and the computation and broadcast of each segment of multipliers for step  $k + 1$ , can be initiated as soon as the relevant segments of row  $k + 1$  and column  $k + 1$  have been updated by their owners, before moving to completing the update for step  $k$
  - ▶ this *send ahead* strategy may grant other tasks earlier access to the data necessary to start working on the next step

# Pivoting

- ▶ the order of rows of  $A$  does not affect the solution to the system of equations
  - ▶ *partial pivoting* sorts the rows by the largest absolute value of the leading column of the remaining unreduced matrix
  - ▶ this choice ensures that the magnitude of the multipliers do not exceed 1, which
    - ▶ reduces amplification of round-off errors
    - ▶ ensures existence
    - ▶ improves numerical stability
- ▶ partial pivoting introduces a permutation matrix  $P$ , which leads to the factorization

$$PA = LU \quad (11)$$

which implies that the solution  $x$  is obtained through

$$Ly = Pb \quad (12)$$

$$Ux = y \quad (13)$$

with forward substitution in the lower triangular system, followed by back substitution in the upper triangular system

# Pivoting in parallel

- ▶ increased numerical stability costs increased parallel complexity and significant performance implications
- ▶ for one dimensional coarsening by column, the search for the pivot element requires no extra communication, but it is purely serial
  - ▶ once the pivot is found, the index of the pivot row must be communicated to the other tasks, and rows must be explicitly or implicitly interchanged in each task
- ▶ for coarsening by rows, the search for the pivot is parallel, but it requires communication among tasks and inhibits the overlapping of successive steps
  - ▶ if rows are explicitly interchanged, then only two tasks are involved
  - ▶ if rows are implicitly interchanged, changes to the assignment of rows to tasks are required, which has effects on concurrency and load balance
- ▶ in the presence of partial pivoting, column and row coarsening trade off on the relative speeds of computation versus communication
- ▶ with two dimensional coarsening, pivot search is parallel but requires communication among tasks along columns and destroys the possibility of overlapping successive steps

# Alternatives to pivoting

- ▶ various alternatives have been proposed
  - ▶ constraining pivoting to blocks of rows
  - ▶ pivoting when the multiplier exceeds a given threshold
  - ▶ pairwise pivoting
- ▶ these strategies are not foolproof, and trade off some stability and accuracy for speed

# Cholesky factorization

- ▶ when  $A$  is a positive definite symmetric matrix it has a Cholesky factorization

$$A = LL^T \quad (14)$$

with  $L$  a lower triangular matrix with positive entries along the diagonal

- ▶ so the linear system  $Ax = b$  can be solved through

$$Ly = b \quad (15)$$

$$L^T x = y \quad (16)$$

- ▶ the factorization is derived by equating corresponding entries of  $A$  with those of  $LL^T$  and generating them in the correct order
  - ▶ for example, in the  $2 \times 2$  case

$$\begin{bmatrix} A_{11} & A_{21} \\ A_{21} & A_{22} \end{bmatrix} = \begin{bmatrix} L_{11} & 0 \\ L_{21} & L_{22} \end{bmatrix} \begin{bmatrix} L_{11} & L_{21} \\ 0 & L_{22} \end{bmatrix} \quad (17)$$

yields

$$L_{11} = \sqrt{A_{11}} \quad L_{21} = A_{21}/L_{11} \quad L_{22} = \sqrt{A_{22} - L_{21}^2} \quad (18)$$



# Computing the Cholesky factorization

---

**Algorithm 6:** CHOLESKY(A)

---

```
1 for  $k = 1$  to  $n$  do
2    $A_{kk} = \sqrt{A_{kk}}$ 
3   for  $i = k + 1$  to  $n$  do
4      $A_{ik} = A_{ik}/A_{kk}$ 
5   for  $j = k + 1$  to  $n$  do
6     for  $i = j$  to  $n$  do
7        $A_{ij} = A_{ij} - A_{ik}A_{jk}$ 
```

---

- note that
  - $n$  square roots are required, all of positive numbers
  - only lower triangle of  $A$  is accessed, so the strict upper triangular part need not be stored
  - $A$  becomes  $L$  in place
  - the algorithm is stable so no pivoting is required
- it takes roughly half the number of  $LU$  operations: approximately  $n^3/6$  multiply-adds

# Parallelizing Cholesky

- ▶ number fine grain tasks as  $(i,j)$  with  $i,j = 1, \dots, n$ ; each task
  - ▶ stores  $A_{ij}$
  - ▶ computes and stores  $L_{ij}$ , if  $i \geq j$
  - ▶ computes and stores  $L_{ji}$ , if  $i < j$

yielding a two dimensional array of  $n^2$  tasks

- ▶ no need to compute and store the zero entries in the upper triangle

# Communication patterns for parallel Cholesky

---

**Algorithm 7:** CHOLESKY( $A$ , task= $(i, j)$ )

---

```
1 for  $k = 1$  to  $\min(i, j) - 1$  do
2   recv  $A_{kj}$ 
3   recv  $A_{ik}$ 
4    $A_{ij} = A_{ij} - A_{ik}A_{kj}$ 
5 if  $i = j$  then
6    $A_{ii} = \sqrt{A_{ii}}$ 
7   broadcast  $A_{il}$  to tasks  $(k, i)$  and  $(i, k)$ ,  $k = i + 1, \dots, n$ 
8 if  $i < j$  then
9   recv  $A_{ii}$ 
10   $A_{ij} = A_{ij}/A_{ii}$ 
11  broadcast  $A_{ij}$  to  $(k, j)$ ,  $k = i + 1, \dots, n$ 
12 if  $i > j$  then
13  recv  $A_{jj}$ 
14   $A_{ij} = A_{ij}/A_{jj}$ 
15  broadcast  $A_{ij}$  to  $(i, k)$ ,  $k = j + 1, \dots, n$ 
```

---

# Coarsening

- ▶ strategies very similar to  $LU$  factorization
  - ▶ one dimensional by row or column
  - ▶ two dimensional blocks

with column coarsening used most often in practice

- ▶ each choice of index in the outer loop yields different algorithm, named after the portion of the matrix that is updated by the basic operation in the inner loops
  - ▶ submatrix Cholesky: with  $k$  as the outer loop index, the inner loops perform a rank 1 update of the remaining unreduced submatrix, using the current column
  - ▶ column Cholesky: with  $j$  in the outer loop, inner loops compute the current column, using matrix-vector multiplies that accumulates the effects of previous columns
  - ▶ row Cholesky: with  $i$  in the outer loop, inner loops compute current row by solving a triangular system involving the previous rows

# Cholesky memory access patterns

