

ACM/CS 114

Parallel algorithms for scientific applications

Michael A. G. Aïvázis

California Institute of Technology

Winter 2010

LU factorization

- ▶ systems of linear equations are ubiquitous in numerical analysis
- ▶ let A be an $n \times n$ matrix, b a known n -vector; we are looking for x such that

$$Ax = b \quad (1)$$

- ▶ a commonly used direct method for solving this system is to convert A into the product of a lower triangular matrix L with an upper triangular matrix U

$$A = LU \quad (2)$$

known as LU factorization

- ▶ Eq. 1 becomes

$$LUx = b \quad (3)$$

which we can now solve in two simpler steps

$$Ly = b \quad (4)$$

$$Ux = y \quad (5)$$

where we first solve the lower triangular system by forward substitution, followed by solving the upper triangular system by back substitution to obtain x

LU by Gaussian elimination

- ▶ we can compute the LU factorization of A using Gaussian elimination

Algorithm 1: LU(A)

```
1 for  $k = 1$  to  $n - 1$  do  
2   for  $i = k + 1$  to  $n$  do  
3      $L_{ik} = A_{ik} / A_{kk}$   
4   for  $j = k + 1$  to  $n$  do  
5     for  $i = k + 1$  to  $n$  do  
6        $A_{ij} = A_{ij} - L_{ik} A_{kj}$ 
```

which encodes L and U in place by overwriting A

- ▶ Alg. 1 requires roughly $n^3/3$ multiply-adds and $n^2/2$ divisions
- ▶ we may also need *pivoting* to ensure numerical stability (and existence)
- ▶ Alg. 1 is one of many algorithms expressed essentially as a triply nested loop
 - ▶ the three indices can be ordered in any of $3!$ ways, with totally different memory access patterns
 - ▶ in parallel, the kji and kji forms may be the most efficient

Parallel LU decomposition

- ▶ number fine grain tasks as (i,j) with $i,j = 1, \dots, n$; each task
 - ▶ stores A_{ij}
 - ▶ computes and stores U_{ij} , if $i \leq j$
 - ▶ computes and stores L_{ij} , if $i > j$

yielding a two dimensional array of n^2 tasks

- ▶ no need to compute and store
 - ▶ the zeroes in the lower triangle of U
 - ▶ the unit diagonal and the zeroes in the upper triangle of L
- ▶ in order to create p coarse grain tasks we could combine
 - ▶ n/p rows or columns of fine grain tasks
 - ▶ $(n/\sqrt{p}) \times (n/\sqrt{p})$ blocks of tasks

and map each one to a process

Communication patterns for parallel LU decomposition

Algorithm 2: $LU(A, \text{task}=(i,j))$

```
1 for  $k = 1$  to  $\min(i,j) - 1$  do
2   recv  $A_{kj}$ 
3   recv  $L_{ik}$ 
4    $A_{ij} = A_{ij} - L_{ik}A_{kj}$ 
5 if  $i \leq j$  then
6   broadcast  $A_{ij}$  to  $(k,j), k = i + 1, \dots, n$ 
7 else
8   recv  $A_{jj}$ 
9    $L_{ij} = A_{ij}/A_{jj}$ 
10  broadcast  $L_{ij}$  to  $(i,k), k = i + 1, \dots, n$ 
```

Row coarsening

- ▶ with one dimensional row coarsening
 - ▶ we forgo parallelism in updating rows
 - ▶ there is no need to broadcast the multipliers L_{ij} since each row is contained entirely within a task
 - ▶ we still need the vertical broadcasts of matrix rows to the tasks below

Algorithm 3: LU(A , task= (i,j)) by rows

```
1 for  $k = 1$  to  $n - 1$  do
2   if  $k \in \text{myrows}$  then
3     broadcast  $\{A_{kj} : k \leq j \leq n\}$ 
4   else
5     recv  $\{A_{kj} : k \leq j \leq n\}$ 
6   for  $i \in \text{myrows}, i > k$  do
7      $L_{ik} = A_{ik}/A_{kk}$ 
8   for  $j = k + 1$  to  $n$  do
9     for  $i \in \text{myrows}, i > k$  do
10       $A_{ij} = A_{ij} - L_{ik}A_{kj}$ 
```

Observations on row coarsening

- ▶ each task becomes idle as soon as its last row is completed
 - ▶ if rows are contiguous, a task may finish long before the overall computation is done
 - ▶ even worse, updating rows requires progressively less work with increasing row number
- ▶ we may improve concurrency and load balance
 - ▶ by assigning rows to tasks in a cyclic manner where row i is updated by task $i \bmod p$
 - ▶ other mappings may be useful
- ▶ other improvements involve overlapping computation with communication
 - ▶ at step k , each task completes updating its portion of the remaining unreduced matrix before moving on to step $k + 1$
 - ▶ however, the task that owns the $k + 1$ row could broadcast it as soon as it becomes available, before moving on to the step k update
 - ▶ this *send ahead* strategy may grant other tasks earlier access to the data necessary to start working on the next step

Column coarsening

Algorithm 4: $\text{LU}(A, \text{task}=(i,j))$ by columns

```
1 for  $k = 1$  to  $n - 1$  do
2   if  $k \in \text{mycolumns}$  then
3     for  $i = k + 1$  to  $n$  do
4        $L_{ik} = A_{ik} / A_{kk}$ 
5       broadcast  $\{L_{ik} : k < i \leq n\}$ 
6   else
7     recv  $\{L_{ik} : k < i \leq n\}$ 
8   for  $i \in \text{mycolumns}, j > k$  do
9     for  $i = k + 1$  to  $n$  do
10       $A_{ij} = A_{ij} - L_{ik}A_{kj}$ 
```

- observations similar to row coarsening apply

Block coarsening

Algorithm 5: LU(A , task= (i, j)) by blocks

```
1 for  $k = 1$  to  $n - 1$  do
2   if  $k \in \text{myrows}$  then
3     broadcast  $\{A_{kj} : j \in \text{mycolumns}, j > k\}$  to all tasks in my task
        column
4   else
5     recv  $\{A_{kj} : j \in \text{mycolumns}, j > k\}$ 
6   if  $k \in \text{mycolumns}$  then
7     for  $i \in \text{myrows}, i > k$  do
8        $L_{ik} = A_{ik} / A_{kk}$ 
9     broadcast  $\{L_{ik} : i \in \text{myrows}, i > k\}$  to all tasks in my task row
10  else
11    recv  $\{L_{ik} : i \in \text{myrows}, i > k\}$ 
12  for  $j \in \text{mycolumns}, j > k$  do
13    for  $i \in \text{myrows}, i > k$  do
14       $A + ij = A_{ij} - L_{ik}A_{kj}$ 
```

Observations on block coarsening

- ▶ each task becomes idle as soon as its last row and column are completed
 - ▶ if rows and columns are in contiguous blocks, a task may finish long before the overall computation is done
 - ▶ even worse, computing multipliers and updating blocks requires progressively less work with increasing row and column numbers
- ▶ we may improve concurrency and load balance
 - ▶ by assigning rows and columns to tasks in a cyclic manner where A_{ij} is assigned to task $(i \bmod \sqrt{p}, j \bmod \sqrt{p})$
 - ▶ other mappings may be useful
- ▶ other improvements involve overlapping computation with communication
 - ▶ at step k , each task completes updating its portion of the remaining unreduced submatrix before moving on to step $k + 1$
 - ▶ the broadcast of each segment of row $k + 1$, and the computation and broadcast of each segment of multipliers for step $k + 1$, can be initiated as soon as the relevant segments of row $k + 1$ and column $k + 1$ have been updated by their owners, before moving to completing the update for step k
 - ▶ this *send ahead* strategy may grant other tasks earlier access to the data necessary to start working on the next step

Pivoting

- ▶ the order of rows of A does not affect the solution to the system of equations
 - ▶ *partial pivoting* sorts the rows by the largest absolute value of the leading column of the remaining unreduced matrix
 - ▶ this choice ensures that the magnitude of the multipliers do not exceed 1, which
 - ▶ reduces amplification of round-off errors
 - ▶ ensures existence
 - ▶ improves numerical stability
- ▶ partial pivoting introduces a permutation matrix P , which leads to the factorization

$$PA = LU \tag{6}$$

which implies that the solution x is obtained through

$$Ly = Pb \tag{7}$$

$$Ux = y \tag{8}$$

with forward substitution in the lower triangular system, followed by back substitution in the upper triangular system

Pivoting in parallel

- ▶ increased numerical stability costs increased parallel complexity and significant performance implications
- ▶ for one dimensional coarsening by column, the search for the pivot element requires no extra communication, but it is purely serial
 - ▶ once the pivot is found, the index of the pivot row must be communicated to the other tasks, and rows must be explicitly or implicitly interchanged in each task
- ▶ for coarsening by rows, the search for the pivot is parallel, but it requires communication among tasks and inhibits the overlapping of successive steps
 - ▶ if rows are explicitly interchanged, then only two tasks are involved
 - ▶ if rows are implicitly interchanged, changes to the assignment of rows to tasks are required, which has effects on concurrency and load balance
- ▶ in the presence of partial pivoting, column and row coarsening trade off on the relative speeds of computation versus communication
- ▶ with two dimensional coarsening, pivot search is parallel but requires communication among tasks along columns and destroys the possibility of overlapping successive steps

Alternatives to pivoting

- ▶ various alternatives have been proposed
 - ▶ constraining pivoting to blocks of rows
 - ▶ pivoting when the multiplier exceeds a given threshold
 - ▶ pairwise pivoting
- ▶ these strategies are not foolproof, and trade off some stability and accuracy for speed

Cholesky factorization

- ▶ when A is a positive definite symmetric matrix it has a Cholesky factorization

$$A = LL^T \quad (9)$$

with L a lower triangular matrix with positive entries along the diagonal

- ▶ so the linear system $Ax = b$ can be solved through

$$Ly = b \quad (10)$$

$$L^T x = y \quad (11)$$

- ▶ the factorization is derived by equating corresponding entries of A with those of LL^T and generating them in the correct order
 - ▶ for example, in the 2×2 case

$$\begin{bmatrix} A_{11} & A_{21} \\ A_{21} & A_{22} \end{bmatrix} = \begin{bmatrix} L_{11} & 0 \\ L_{21} & L_{22} \end{bmatrix} \begin{bmatrix} L_{11} & L_{21} \\ 0 & L_{22} \end{bmatrix} \quad (12)$$

yields

$$L_{11} = \sqrt{A_{11}} \quad L_{21} = A_{21}/L_{11} \quad L_{22} = \sqrt{A_{22} - L_{21}^2} \quad (13)$$

Computing the Cholesky factorization

Algorithm 6: CHOLESKY(A)

```
1 for  $k = 1$  to  $n$  do
2    $A_{kk} = \sqrt{A_{kk}}$ 
3   for  $i = k + 1$  to  $n$  do
4      $A_{ik} = A_{ik}/A_{kk}$ 
5   for  $j = k + 1$  to  $n$  do
6     for  $i = j$  to  $n$  do
7        $A_{ij} = A_{ij} - A_{ik}A_{jk}$ 
```

- ▶ note that
 - ▶ n square roots are required, all of positive numbers
 - ▶ only lower triangle of A is accessed, so the strict upper triangular part need not be stored
 - ▶ A becomes L in place
 - ▶ the algorithm is stable so no pivoting is required
- ▶ it takes roughly half the number of LU operations: approximately $n^3/6$ multiply-adds

Parallelizing Cholesky

- ▶ number fine grain tasks as (i,j) with $i,j = 1, \dots, n$; each task
 - ▶ stores A_{ij}
 - ▶ computes and stores L_{ij} , if $i \geq j$
 - ▶ computes and stores L_{ji} , if $i < j$

yielding a two dimensional array of n^2 tasks

- ▶ no need to compute and store the zero entries in the upper triangle

Communication patterns for parallel Cholesky

Algorithm 7: CHOLESKY(A , task= (i, j))

```
1 for  $k = 1$  to  $\min(i, j) - 1$  do
2   recv  $A_{kj}$ 
3   recv  $A_{ik}$ 
4    $A_{ij} = A_{ij} - A_{ik}A_{kj}$ 
5 if  $i = j$  then
6    $A_{ii} = \sqrt{A_{ii}}$ 
7   broadcast  $A_{il}$  to tasks  $(k, i)$  and  $(i, k)$ ,  $k = i + 1, \dots, n$ 
8 if  $i < j$  then
9   recv  $A_{ii}$ 
10   $A_{ij} = A_{ij}/A_{ii}$ 
11  broadcast  $A_{ij}$  to  $(k, j)$ ,  $k = i + 1, \dots, n$ 
12 if  $i > j$  then
13  recv  $A_{jj}$ 
14   $A_{ij} = A_{ij}/A_{jj}$ 
15  broadcast  $A_{ij}$  to  $(i, k)$ ,  $k = j + 1, \dots, n$ 
```

Coarsening

- ▶ strategies very similar to LU factorization
 - ▶ one dimensional by row or column
 - ▶ two dimensional blocks

with column coarsening used most often in practice

- ▶ each choice of index in the outer loop yields different algorithm, named after the portion of the matrix that is updated by the basic operation in the inner loops
 - ▶ submatrix Cholesky: with k as the outer loop index, the inner loops perform a rank 1 update of the remaining unreduced submatrix, using the current column
 - ▶ column Cholesky: with j in the outer loop, inner loops compute the current column, using matrix-vector multiplies that accumulates the effects of previous columns
 - ▶ row Cholesky: with i in the outer loop, inner loops compute current row by solving a triangular system involving the previous rows

Cholesky memory access patterns

